

ragfarm — on-prem RAG + infra-control agent

David Kubicek (kubicek@gmail.com)

August 14, 2026

Contents

Chapter 1 Introduction	7
Introduction — what this document is and why it exists	7
Why on-prem RAG matters right now	7
What is genuinely custom about this build	8
What this bundle contains and how to read it	9
Chapter 2 Project README	11
ragfarm — on-prem RAG + infra-control agent	11
Hardware	11
The one fact that drives every decision	11
Engine split (ADR-0013)	12
Quick start	12
Retrieval pipeline	12
Layout	12
Diagrams	13
Runtime component inventory	13
Query-time control flow	14
Systemd unit topology (host services)	15
Container network topology (infra/compose.yaml)	16
Working chat examples (verified in OWUI)	16
1. Firewall rules for a specific host (RAG → structured table)	16
2. Contacts for the EPC project team (RAG → per-person structured ...	17
3. Deep host lookup — 19 fields for one server (RAG → keyed detail ...	18
4. Where credentials live (RAG + procedural answer)	19
5. Documentation Git repository (RAG → exact URL)	20
6. Access flow from ŠA into EPC (RAG → 6-way procedural breakdown)	20
7. Tool discipline: time query, out-of-corpus refusal, gated reboot	21

8. Diagram generation — dependency tree from a natural-language . .	22
9. Code generation with in-chat execution	24
Network / proxy	24
Where to start	25
Status	25
Chapter 3 Qwen3-VL vision demonstrations	26
Qwen3-VL demonstrations — reproducible board-demo examples	26
1 · RAG over internal corpus — one tool call, full-column table	26
2 · Image classification — "what's on this picture?"	27
3 · OCR of a hand-written German book page	27
4 · draw.io round-trip — reverse every arrow direction	28
5 · Mermaid generation for a stack diagram (natural-language spec)	29
6 · Where the current iGPU hardware hits a wall	30
What this chapter demonstrates	30
Chapter 4 Hardware requirements (dev + prod ladder)	32
Hardware requirements — how they scale, and what to buy	32
Tier 0 — the original PoC box (HISTORICAL — kept as the bottom rung) .	32
The one formula — how to size any of this	33
The constant, back-solved from our own hardware	34
The first consequence: never buy dense	34
Sizing a target	34
Tier 0.5 — NVIDIA DGX Spark (the current substrate)	35
Tier 1 — Developer workstation	36
Tier 2 — Small-team production server (10-50 concurrent users)	37
Tier 3 — DC-tier serious production (100s of concurrent users, or a	38
The trajectory that matters for planning	39
The concrete ask, if you have one conversation to make it in	40
Chapter 5 Deployment guide	42
Deployment & prod re-deployment notes	42
Runtime topology — live component registry	42
Why host networking (load-bearing PoC fact)	44
Remote UI access	45
Tested model versions (known-good baseline)	45
Autostart & lifecycle	46
What starts on boot	46
Start / stop / status everything (host, as dave)	46
Restarting individual pieces (the gotchas)	47
GPU contention & the reranker's priority (tunable — revisit with	47
Operator scripts (scripts/) — a newcomer's map	48

Open WebUI configuration (reproducible, not hand-clicked)	51
Verifying the toolchain	52
OpenNebula MCPs — mock mode & the confirmation gate (ADR-0004) . .	52
The AMD → Spark migration (DONE) and what is still open	53
Vision engine (Qwen3-VL family — ADR-0009)	54
Live capabilities that Just Work	54
Diagram rendering	55
draw.io: the six things that must all be true	55
Verified demo commands (no prep needed, run today)	57
/think vs /no_think (Qwen3 Thinking control)	58
Debug & measurement (tests/tracing/)	58
The catalog	59
Verified one-liners (run today)	60
What "good" looks like — Spark (current) vs AMD iGPU (historical) . . .	60
Slots — two models resident, switchable mid-chat	62
Chapter 6 LLM life-cycle — models, slots, running the stack	63
LLM life-cycle — models, slots, and running the stack	63
What a slot is	63
The memory budget, which is the whole difficulty	63
The registry	64
Everyday commands	64
Four things that will bite	65
Boot behaviour	65
Manual pages	66
Chapter 7 Prompt library and regression suite	67
Prompt library — demo reference and regression suite	67
How to add a case	67
Section R — Retrieval and grounding	68
R1 · FW rules for one specific host	68
R2 · Contacts for project management	68
R3 · Everything about one host	69
R4 · Where credentials live	69
R5 · Where the documentation repo is	69
R6 · How to log in from ŠA to EPC	69
R7 · All CutOver contacts	70
R8 · General knowledge is answered directly	70
R9 · Honest miss inside the corpus domain	70
Section T — Tools beyond retrieval	70
T1 · Clock	70

T2 · Human-gated reboot	71
Section V — Vision	71
V1 · Receipt OCR, real world	71
V2 · Mixed-script OCR with identifiers	71
V3 · Chinese to English	72
V4 · Farsi to English	72
V5 · Bar chart to table	72
V6 · Diagram photo to mermaid	72
V7 · Code photo to runnable Python	72
Section D — Diagrams the model authors	73
D1 · Reverse the arrows in a supplied diagram	73
D2 · Author a diagram from a description	73
D3 · 1:1 conversion of a diagram image	73
D4 · Inventory a diagram image without drawing it	74
Section C — Coding	74
C1 · Write and actually run Python	74
C2 · Non-Python code is not executed	74
Section M · Mermaid	75
M1 · Dependency tree of a sentence	75
Chapter 8 Prompt regression testing	76
Prompt regression testing	76
Why this exists	76
The library is the suite	76
Running it	76
Two phases, and why they are separate	77
Verdicts	77
Calibrating the judge	77
What is not covered yet	78
See also	78
Chapter 9 Ingestion pipeline	79
Ingestion pipeline — mixed docx / xlsx / pdf corpus	79
1. File routing	79
2. Table path (xlsx / csv / docx-tables) — one row per chunk	79
Messy real-world XLSX	80
3. Prose path (docx / pdf / txt / md) — section-aware chunks	80
4. Embedding + storage	81
Collection schema + alias (created on first run)	82
5. Sync, idempotency & re-runs (ADR-0006 — content-addressed)	82
6. Why hybrid (dense + sparse) for this corpus specifically	83

7. The automatic watcher (deployed autonomous sync)	83
Deployment	83
How a pass is triggered	83
What a pass does	84
Concurrency & robustness	84
Deployed knobs (unit values; all env-overridable)	84
8. Verification (feeds step 04's gate)	85
Chapter 10a Embedder README	86
Embedder — BAAI/bge-m3 on CPU (:8090/embed)	86
One endpoint, one job	86
Contract	86
Run	86
Chapter 10b llama.cpp (reranker engine)	87
llama.cpp — CUDA build, for the reranker only	87
Build	87
Running it	88
See also	88
Chapter 11 Configuration reference (.env.example)	89
Chapter 12a Model record — LLM	98
Generative LLM Model Record — the Qwen3-VL fleet	98
Baseline record — Qwen3-VL-30B-A3B-Instruct (NVFP4)	99
MoE backend on sm_121 — MEASURED, and it contradicts ADR-0013 . . .	99
Startup memory — the two budgets are NOT the same lever	100
Verified	101
NOT yet measured	101
Chapter 12b Model record — embedder	102
Embedding Model Record	102
Chapter 12c Model record — reranker	103
Reranker Model Record (ADR-0008)	103
Why a second llama.cpp server (not an embedder sub-endpoint)	103
Regenerating the GGUF (weights are gitignored)	103
Chapter 13 Instrumentation and tracing	105
Instrumentation and tracing — what these tools are, and what they	105
What the tracing work established	105
What is wrong with them today	105
Endpoint resolution	106
The tools	106
What superseded them	107
Requirements for the rewrite	107

Chapter 14 Build progress (PROGRESS.md)	108
PROGRESS — blocker channel between the agent and Dave	108
Format	108
Entries	108

Chapter 1 Introduction

Introduction — what this document is and why it exists

This bundle is the single authoritative reference for **ragfarm**: a fully on-prem retrieval-augmented generation (RAG) assistant with infrastructure-control capabilities, running today on an **NVIDIA DGX Spark (GB10 Grace Blackwell)** and portable, unchanged, to larger NVIDIA hardware.

It is generated on demand from the repository's live source-of-truth documents. No content exists only here — every chapter has a stable home in the repo tree, and this bundle concatenates and typesets them into one printable, offline artifact. When a source document changes, `docs/pdf/build.sh` regenerates the PDF in seconds.

A note on the historical material. The project began on an AMD Ryzen AI mini-PC and moved to the Spark in August 2026. Where old numbers appear — 8 tokens/second decode, three-minute Farsi OCR, the hardware ladder's Tier 0 — they are kept deliberately and labelled. They are the bottom rung, and the only tier where we have lived with the consequences of being under-provisioned.

Why on-prem RAG matters right now

Two forces meet at this moment. The first is the practical maturity of open-weight transformer models in the 8-30 billion-parameter range — specifically the **mixture-of-experts** generation, where a 30-billion-parameter model activates only about 3 billion per token. These are now competent at structured document retrieval, tabular extraction, tool-driven infrastructure operations, and, in the vision variants, direct optical understanding of scanned pages, screenshots and diagrams. They fit on hardware an operator can buy once.

The second is the regulatory and commercial tension around cloud AI. Corporate documentation, employee contacts, network topology, credentials and operational runbooks either cannot legally leave the enterprise perimeter (energy, finance, healthcare, defence, NIS2) or represent competitive information whose loss to a third-party provider is unacceptable. Cloud AI trades ease of setup for a permanent flow of internal data through someone else's servers, a permanent per-token cost, and lock-in to whichever vendor best matches this quarter's price/quality trade-off.

ragfarm demonstrates that the trade is optional. The same natural-language interface — the same OpenAI-compatible chat, the same tool calling, the same vision — runs entirely inside the perimeter, with no per-token cost, no data egress, and no vendor dependency for the working system.

What is genuinely custom about this build

Most open-source LLM stacks are upstream projects glued together with wiring code. This one is honestly the same, and the wiring is deliberately thin so the pieces stay independently upgradeable. Six areas produced measurable improvement over the out-of-the-box configuration. These are the parts worth reproducing.

1. The engine split, decided by measurement.

Generation runs on **vLLM** with an NVFP4 mixture-of-experts checkpoint; the cross-encoder reranker stays on **llama.cpp** at lowered scheduling priority so the interactive model wins contention; the BGE-M3 embedder runs on CUDA; Qdrant and the MCP services run on CPU. All three GPU consumers share one memory pipe, which is why every model choice is argued against bandwidth rather than parameter count (ADR-0013).

The sm_121 FP4 backend choice deserves its own warning: a misconfigured NVFP4 MoE on GB10 fails **silently** — the model loads cleanly and then emits streams of !!!!!. Garbage output here is a kernel problem first and a model problem second.

2. A measured hardware-sizing model, not a vendor slide.

Decode is memory-bandwidth-bound, so tokens/second is bandwidth divided by bytes read per token, and bytes read follows active parameters. We back-solved the constant from our own two models: a 30B-A3B MoE at 68 tok/s and a dense 32B at 5.9 tok/s both imply **200 GB/s effective, 73% of specification**. The same size class, **11.5× apart**. That single ratio is the argument for buying bandwidth and MoE rather than parameter count, and Chapter 4 turns it into priced options.

3. Retrieval tuned from stock BGE-M3 to a domain-adapted

hybrid. Dense (1024-dim) plus native sparse, fused with Reciprocal Rank Fusion, re-ranked by a bge-reranker-v2-m3 cross-encoder, then cut by a floor plus a Kneedle chord-distance gate. The reranker is the most expensive stage at 297 ms and decides whether the model receives the right passages at all — which is why it is never shrunk to buy latency. Also custom: broad-in/narrow-out chunking tuned for structured XLSX and Confluence-style prose, and an XLSX table parser that recovers headers from multi-row banded layouts where standard tools return zero rows (ADR-0007, ADR-0008, ADR-0010).

4. The grounding prompt as a seven-rule contract — and the

discovery that it is the whole specification. RULE 0 through RULE 6 in `infra/openwebui/setup_openwebui.py` define when a tool is required, how much to retrieve, when a table is the answer and when one record is, how images are handled, which diagram format to emit, and that only Python is executed.

The important finding is methodological. Every "the model is stupid" symptom investigated in August turned out to be **an instruction defect, not a model defect**: a rule saying "call the tool once, do not refine" was generalised into be minimal, so the model set $k=1$ and starved its own retrieval. Fixing the prompt took a model from 2/5 to 9/10 on grounded questions and shrank its reasoning from 16,805 to 2,788 characters. That is a larger gain than any hardware upgrade on this page, and it cost nothing.

Because the prompt is that load-bearing, it now has a **regression suite** (Chapters 7 and 8): the prompt library is machine-readable, replayed against the live model, and judged by a second model call.

5. Slots — two models resident at once, switchable

mid-conversation. vLLM serves one model per process, so two resident models means two instances with a shared, non-coordinating memory budget. The tooling derives each slot's share from the verified checkpoint size and refuses any binding that would exceed the ceiling. The payoff is diagnostic: asking a second model the same question with the same history separates "the model cannot do this" from "our prompt told it not to" (Chapter 6).

6. draw.io rendered in-chat, air-gap-safe.

A vision model can regenerate a scanned architecture diagram as mermaid or as an interactive draw.io canvas, served from a full local mirror so nothing leaves the box at demo time. Getting there required fixing four independent faults that all presented as the same blank white rectangle — a missing mirror, a loopback URL that a remote browser resolves to itself, a content-security policy blocking the corrected URL, and an XML hand-off form the viewer does not read. Every one of them failed **silently**. That story is in the deployment guide, and it is the best illustration in this bundle of why "it returned 200" is not evidence of anything.

What this bundle contains and how to read it

chapters	what
1-4	This introduction, the README, vision demonstrations, and the hardware ladder with the sizing formula

5-6	Deployment guide, then the LLM life-cycle — models, slots, and running the stack
7-8	The prompt library, then how it doubles as a regression suite
9-10	Ingestion pipeline and component READMEs
11-12	Configuration reference and model records
13	Instrumentation and tracing: the tools, and what they taught us
14	Build progress — the operational ledger on the day this PDF was built

For a technical audience: start with the deployment guide and the life-cycle chapter, and treat the rest as backing material.

For a business audience: this introduction plus Chapter 4 (hardware) are sufficient, and Chapter 4 is where the money question is answered.

For operators inheriting the system: read every chapter once, then keep the deployment guide and `man docs/man1/stack.1` open.

Chapter 2 Project README

ragfarm — on-prem RAG + infra-control agent

Author: David Kubicek (david.kubicek@eywo.cz) · MIT, see LICENSE.txt

An on-prem retrieval-augmented assistant over a customer VM farm. It answers questions about thousands of internal documents **and** takes operational requests in natural language ("where is VM1 running?", "reboot host X"), which are dispatched through MCP microservices to OpenNebula.

Nothing leaves the box. No per-token cost, no data egress, no vendor dependency for the working system.

Hardware

Box	NVIDIA DGX Spark — GB10 Grace Blackwell, sm_121
Memory	128 GB unified (121.7 GiB usable), CPU and GPU share it
Bandwidth	273 GB/s spec · ** 200 GB/s measured effective**
OS	Ubuntu 24.04, Python 3.12

Previously an AMD Ryzen AI 9 HX 370 mini-PC; that hardware is retired and its numbers are kept only as the bottom rung of the ladder in docs/pdf/02-hardware-requirements.md.

The one fact that drives every decision

Decode is memory-bandwidth-bound. Tokens per second is bandwidth divided by bytes read per token, and bytes read follows **active** parameters, not total.

Measured here, two models of the same size class:

model	active bytes/token	tok/s	implied bandwidth
Qwen3-VL-30B-A3B FP8 — MoE, 3 B active	3 GB	68	204 GB/s
Qwen3-VL-32B FP8 — dense, 32 B active	33 GB	5.9	195 GB/s

11.5x apart. Hence: mixture-of-experts always, dense never, and every "just add another model" argued against the memory budget rather than against disk. Sizing formulas and priced hardware options are in chapter 4 of the PDF bundle.

Engine split (ADR-0013)

component	runs on	endpoint
Generative LLM	vLLM, CUDA, NVFP4/FP8	127.0.0.1:8080 (+ :8082 for slot 1)
Reranker bge-reranker-v2-m3	llama.cpp, CUDA, low-ered priority	127.0.0.1:8081/reranking
Embedder BGE-M3 dense+sparse	CUDA	127.0.0.1:8090/embed
Qdrant + MCP services	CPU	:6333, :8000, :8101-8104

Two vLLM **slots** can hold two models at once, switchable mid-conversation with the context intact — see docs/llm-lifecycle.md.

sm_121 warning. A misconfigured NVFP4 MoE on GB10 fails silently: the model loads and then emits streams of !!!!!. Garbage output is a backend/kernel problem first, a model problem second.

Quick start

```
scripts/stack.sh status      # 13 services: endpoint, state, depth checks
scripts/stack.sh start      # host services, then containers, then health
scripts/activate_llm.py --status      # slots, models, GPU budget
scripts/activate_llm.py      # interactive: activate or clear
```

Open WebUI: `http://<host>:3000` — the only LAN-exposed service, login-gated.

Manual pages: man docs/
man1/{stack,activate_llm,fetch_llm,active.json,setup_openwebui,env}.1

Retrieval pipeline

BGE-M3 dense (1024-dim) + native sparse → Qdrant two-branch prefetch → RRF fusion → cross-encoder rerank → floor + Kneedle chord-distance gate → top-k → same-section expansion. Measured per stage:

embed 183 ms · fuse 15 ms · rerank 297 ms · expand 6 ms

The reranker is the expensive stage and decides whether the model sees the right passages at all, so it is never shrunk to buy latency. On the old CPU-only placement the same stage took **36 seconds**.

Layout

path	contents
------	----------


docs/	architecture decisions, deployment, prompt library, man pages, measurements
infra/	compose stack, Open WebUI config, embedder, llama.cpp (reranker)
models/	checkpoints + active.json registry (weights are not repo-synced)
services/	ingester, RAG pipeline, MCP microservices
manifests/	systemd units
scripts/	life-cycle tools, regression suite, benchmarks
tests/	fixtures, prompt library parser, tracing tools

Diagrams

Runtime component inventory

ragfarm — runtime component inventory


Every process alive at demo time, grouped by execution plane. Boxes are colored by where each runs; edges are the tools/APIs each speaks.



Query-time control flow

ragfarm — query-time control flow


One user question, top-to-bottom. Yellow numbered chevrons mark the sequence; per-stage timings are current iGPU measurements (see docs/deployment.md → Debug & measurement).



Systemd unit topology (host services)

ragfarm — systemd unit topology (host services)


Every unit installed by `scripts/deploy.sh`, its dependencies, restart behavior, and the ports it binds. Docker itself is a host unit (`docker.service`) that hosts the container stack; see the container diagram for what's inside.



Container network topology (infra/compose.yaml)

ragfarm — container network topology (infra/compose.yaml)

What runs in Docker, on which host port, and which containers use `network_mode: host` to reach loopback services on the box.



Working chat examples (verified in OWUI)

Nine screenshots from the deployed Open WebUI. Every question was answered live, unedited.

Provenance, because it matters for what these prove. These were captured 2026-07-27 on the retired text preset (Qwen2.5-7B, greedy) on the AMD box. They are evidence that the pipeline does what the docs describe, and they are the source of the R1-R6 cases in docs/prompts.md — but they are **not** a record of the current models. Re-running them against today's stack is what scripts/test_regressions.py is for. The same caveat applies to the systemd diagram above, which still shows ragfarm-llama where the deployment now runs ragfarm-vllm@N.

1. Firewall rules for a specific host (RAG → structured table)

Prompt: Jaká jsou FW pravidla pro host leadb229p.lea.piz?

Jaká jsou FW pravidla pro host leadb229p.lea.piz?

oi ragfarm (corpus RAG + infra)

🟢 View Result from tool_search_corpus_post ▾

Zde jsou firewall pravidla pro host leadb229p.lea.piz:

SOURCE NETWORK ADDRESS(ES)	SOURCE NETWORK NAME(S)	DESTINATION NETWORK ADDRESS(ES)	DESTINATION NETWORK NAME(S)	DESTINATION PORT(S)	PROTOCOL
10.244.1.26	leadb229p.lea.piz	172.20.220.0/24	EPC_AZURE	445	TCP
10.244.0.26	leadb229p.lea.piz	172.20.231.116	prdwcamel.epcomm.cz	445	TCP
10.244.0.26	leadb229p.lea.piz	172.20.231.119	prdlsto.epcomm.cz	1521	TCP
172.20.231.116	prdwcamel.epcomm.cz	10.244.0.26	leadb229p.lea.piz	445	TCP
172.20.231.119	prdlsto.epcomm.cz	10.244.0.26	leadb229p.lea.piz	1521	TCP
172.20.220.0/24	EPC_AZURE	10.244.1.26	leadb229p.lea.piz	445	TCP

Tyto pravidla byly požadovány dne 13.7.2026 a jsou založeny na požadavcích Petr Pyszko z EPCModities.

Model calls `search_corpus`, gets 6 firewall-rule rows back, renders them as a markdown table with every column (Source/Destination Network Address(es) and Name(s), Destination Port(s), Protocol) — no dropped fields. Ends with a one-line source citation identifying the .xlsx it came from.

2. Contacts for the EPC project team (RAG → per-person structured

list) **Prompt:** Dej mi kontakty na projektové vedení v EPC.

Dej mi kontakty na projektové vedení v EPC.

oi ragfarm (corpus RAG + infra)

✓ View Result from `tool_search_corpus_post` ▾

Zde jsou kontakty pro projektové vedení v EPC:

1. **Petr Pyzsko**
 - Firma: EPC
 - Role/oblast: Sítě
 - Tel: 737924492
 - E-mail: petr.pyzsko@epcommodities.cz
2. **Miroslav Laboj**
 - Firma: EPC
 - Role/oblast: Aplikace a sítě
 - E-mail: miroslav.laboj@epcommodities.cz
3. **Viktor Bobro**
 - Firma: EPC
 - Role/oblast: Aplikace
 - E-mail: viktor.bobro@epcommodities.cz
4. **Michal Sterba**
 - Firma: EPC
 - Role/oblast: Aplikace
 - E-mail: michal.sterba@epcommodities.cz
5. **Viktor Vážný**
 - Firma: EPC
 - Role/oblast: IT
 - Tel: 604 795 022
 - E-mail: vit.vazny@epcommodities.cz
6. **Marek Česal**
 - Firma: EPC
 - Role/oblast: PM - řízení cut-over procesu
 - Tel: 739223474
 - E-mail: marek.cesal@epcommodities.cz

Pokud potřebujete další informace nebo kontakty, neváhejte se zeptat.

Six people, each rendered as Firma / Role/oblast / Tel / E-mail. Phone numbers appear where present in source; the two people whose corpus row lacks a phone number simply don't get the field (no hallucination).

3. Deep host lookup — 19 fields for one server (RAG → keyed detail view) **Prompt:** Co vis o hostu acclass1?

Co vis o hostu acclclass1?

oi ragfarm (corpus RAG + infra)

🟢 View Result from tool_search_corpus_post ▾

Host **acclclass1** má následující specifikace:

- **Prostředí:** ACC/UAT
- **OS:** Red Hat 8
- **Virtualizace:** KVM
- **Storage:** Huawei
- **vCPU:** 8
- **RAM (GB):** 32
- **HDD (TB):** 1
- **Support:** 7/12
- **T-S Hostname:** DFLEA237
- **T-S IP address:** 10.123.142.70
- **T-S Netmask:** /27
- **SA Hostname:** acclclass1
- **SA domain suffix:** epcomm.cz
- **SA IP address:** 172.20.231.158
- **SA Netmask:** /27
- **SA VLAN ID:** 605
- **Storage OS [GiB]:** 64
- **Storage App/Data [GiB]:** 300
- **App/Data Directories/Filesystems:** /global
- **UID:** aol:users (16573)

Pokud potřebujete další informace nebo konkrétní detaily, neváhejte se zeptat.

Prostředí, OS, Virtualizace, Storage, vCPU, RAM, HDD, Support, T-S Hostname/IP/Netmask, SA Hostname/domain/IP/Netmask/VLAN ID, Storage OS+App/Data, filesystem, UID — all recovered from the source spreadsheet's row for this host, verbatim.

4. Where credentials live (RAG + procedural answer)

Prompt: Kde ukládáme hesla pro EPC?

Kde ukládáme hesla pro EPC?

oi ragfarm (corpus RAG + infra)

View Result from tool_search_corpus_post

Hesla pro EPC jsou obvykle uložena v NordPass. Specifickým krokem je:

1. Přihlaste se k NordPass použitím svého uživatelského jména a hesla pro IdM/IPA.
2. V NordPass najdete záznam s názvem "RDP User" a zapisujete si zde heslo.

Tento záznam obsahuje heslo, které je potřeba zadat v RDP klientu při připojení k EPC terminálovým serveru.

Pokud potřebujete další podrobnosti nebo potřebujete přístup k NordPass, neváhejte se zeptat.

Answers where (NordPass), which record name to look for ("RDP User"), and how that password fits into the RDP-to-terminal-server workflow.

5. Documentation Git repository (RAG → exact URL)

Prompt: Kde máme uložené GIT repo s dokumentací?

Kde máme uložené GIT repo s dokumentací?

oi ragfarm (corpus RAG + infra)

View Result from tool_search_corpus_post

GIT repo s dokumentací je uložený na adrese:

1 https://SGC-DevOps@dev.azure.com/SGC-DevOps/Internal/_git/sa-hosting

⌵ Collapse Save Copy

Tento repo obsahuje dokumentaci a další soubory související s ŠA hostingem. Pokud potřebujete další informace nebo přístup k tomuto repo, neváhejte se zeptat.

Exact repository URL retrieved verbatim from the corpus.

6. Access flow from ŠA into EPC (RAG → 6-way procedural breakdown)

Prompt: Jak se přihlásím ze ŠA do EPC?

oi ragfarm (corpus RAG + infra) 07/20/2026 at 3:41 PM

View Result from tool_search_corpus_post

Z ŠA můžete přihlásit do EPC následujícími způsoby:

1. SSH přímo na hostity:

- SSH na `hmbvxi001ts.csr01.hs.skoda.vwg` (IP: 100.80.9.11) nebo `hmbvxi002ts.csr01.hs.skoda.vwg` (IP: 100.80.9.12) uživatelem v Hosting IdM (např. `mstrnad`).
- Použití uživatelského jména `gnome-remote-desktop` a hesla z NordPass jako "RDP User". Toto vás připojí k přihlašovací obrazovce Gnome. Zde pokračujte svým uživatelem v Hosting IdM.

2. RDP přímo na hostity:

- RDP na tytéž adresy (100.80.9.11 nebo 100.80.9.12) uživatelem v Hosting IdM.

3. Přes terminálové servery v EPC tenantu:

- SSH na `epcns1.epcomm.cz` (IP: 172.20.231.41) nebo `epcns2.epcomm.cz` (IP: 172.20.231.42) uživatelem `hostadmin` (root ho má nastaveného v `.ssh/config`).

4. GUI přímo na EPC terminálové servery:

- Použití URL `http://100.80.9.251:2616/fireedge/sunstone/` s přihlašovacími údaji IdM/IPA.

5. Přes reverse proxy EPC terminálové servery:

- Použití URL `http://epcproxy.epcomm.cz:8080/fireedge` nebo `http://epcns1.epcomm.cz:8080/fireedge` nebo `http://epcns2.epcomm.cz:8080/fireedge`.

6. CLI přímo na hostity:

- Nainstalování CLI klienta na `hmbvxi001ts/002ts`.
- Nastavení proměnné prostředí `ONE_XMLRPC=http://100.80.9.251:2633/RPC2`.
- Vytvoření adresáře `~/one` a nastavení pravidel.
- Zapsání přihlašovacích údajů ve formátu `<IPA user>:<IPA heslo>` do souboru `~/one/one_auth`.


Pro další informace o architektuře EPC VMek, viz soubor `EPC25_VMs_config.xlsx` a pro přehled infrastruktury VMek, viz soubor `SA_Hosting_infra_VMs.xlsx`.

Pokud máte další otázky nebo potřebujete další podporu, neváhejte se zeptat.

Six numbered login methods (SSH direct, RDP direct, via terminal servers, GUI, reverse proxy, CLI), each with the exact hostnames, IPs, credentials source, and pointers back to the source spreadsheets (`EPC25_VMs_config.xlsx`, `SA_Hosting_infra_VMs.xlsx`).

7. Tool discipline: time query, out-of-corpus refusal, gated reboot

Prompts (multi-turn): Rebootuj host node-03. • Kolik je hodin? • Kdo je prezident USA? • Rebootuj host node-03.



Four turns in one image. Reboot fires the `reboot_host` tool (returns success with live-migration report of `sftp-gw`). Time query fires `get_current_timestamp`. **"Who is the president of the USA?"** correctly returns a graceful miss (v dokumentu není zmíněn žádný prezident USA) instead of hallucinating a name — this is the target behavior for anything outside the corpus. Fourth turn reboots the same host again, cleanly.


8. Diagram generation — dependency tree from a natural-language

sentence **Prompt:** Vygeneruj stromový diagram slovních vazeb ve větě: "Once upon a time there was a very little dog called Steven who owned a nice little yellow car".

Vygeneruj stromový diagram slovních vazeb ve větě: "Once upon a time there was a very little dog called Steven who owned a nice little yellow car".

oi ragfarm (corpus RAG + infra)

Zde je Mermaid diagram reprezentující stromový diagram slovních vazeb ve větě: "Once upon a time there was a very little dog called Steven who owned a nice little yellow car."



Mermaid syntax rendered inline in-chat by OWUI as an interactive SVG. Notice this is a real dependency tree, not a linear word chain: *dog* is a hub with *a/very/little/called Steven* branching off; *owned* hangs from *who*; the second *a* is its own hub for *car*. This shape is what took several rounds of system-prompt tuning to reach reliably.

9. Code generation with in-chat execution

Prompt: Vygeneruj mi kód pro quicksort a otestuj ho spuštěním nad malým polem náhodných řetězců. Prezntuj kód a výsledné pořadí tříděného pole po běhu sortu.

Vygeneruj mi kód pro quicksort a otestuj ho spuštěním nad malým polem náhodných řetězců. Prezntuj kód a výsledné pořadí tříděného pole po běhu sortu.

oi ragfarm (corpus RAG + infra)

View Result from execute_code

Zde je generovaný kód pro algoritmus quicksort a výsledek jeho spuštění nad polem náhodných řetězců:

```
python
1 import random
2 import string
3
4 # QuickSort algorithm
5 def quicksort(arr):
6     if len(arr) <= 1:
7         return arr
8     else:
9         pivot = arr[len(arr) // 2]
10        left = [x for x in arr if x < pivot]
11        middle = [x for x in arr if x == pivot]
12        right = [x for x in arr if x > pivot]
13        return quicksort(left) + middle + quicksort(right)
14
15 # Generate a list of 10 random strings
16 random_strings = [''.join(random.choices(string.ascii_letters + string.digits, k=10)) for _ in range(10)]
17
18 # Shuffle the list to get a random order
19 random.shuffle(random_strings)
20
21 # Sort the list using quicksort
22 sorted_strings = quicksort(random_strings)
23
24 sorted_strings
```

Po spuštění sortu je pole tříděné do následujícího pořadí:

TEXT

7BSNAKfI4w

Bqznld4jNr

NaCN7znRml5

Model writes a complete quicksort in Python, OWUI's built-in code interpreter runs it, and the sorted output appears inline as a data grid. Same sampling-deterministic path is used for RULE 5 in the system prompt (write → execute → report + benchmark).

Network / proxy

Outbound build traffic honours a proxy from repo-root .env (gitignored; copy .env.example). **Source the loader before any networked command:**

`source scripts/proxy-env.sh`

`docker compose -f infra/compose.yaml up -d`

The loader merges an internal-host baseline into `NO_PROXY` (`localhost,127.0.0.1,::1,host.docker.internal,qdrant`), so the stack's many loop-back calls never route at the proxy. List only additional bypass targets.

Container image **pulls** go through the Docker daemon, not these variables — configure `/etc/systemd/system/docker.service.d/http-proxy.conf` separately.

Full propagation rules: `man docs/man1/env.1`.

Where to start

you are	read
operating it	<code>docs/deployment.md</code> , then <code>man docs/man1/stack.1</code>
changing models	<code>docs/llm-lifecycle.md</code>
changing the prompt	<code>docs/regression-testing.md</code> — run the suite before and after
evaluating the design	<code>docs/decisions/ADR-0013</code> (engine split), <code>ADR-0003</code> (agent layer), <code>ADR-0010</code> (retrieval gate)
buying hardware	chapter 4 of the PDF bundle

`CLAUDE.md`, `BUILD_STATE.md` and `PROGRESS.md` are the build contract and ledger — written for an AI agent executing the build, of limited use to a human except as a record of what was done and what is blocked.

Status

All architectural decisions through `ADR-0013` are implemented. Two 30B-class MoE models run resident and switchable. Retrieval is tuned for infrastructure documentation — hybrid sparse+dense, RRF, cross-encoder rerank, Kneedle gate. The prompt library doubles as a regression suite.

Known gaps are listed in `PROGRESS.md`; the honest ones worth knowing are that image-input cases are not yet covered by the regression suite, and the tracing tools predate thinking models.

Chapter 3 Qwen3-VL vision demonstrations

Qwen3-VL demonstrations — reproducible board-demo examples

Every example in this chapter was fired against the live stack on 2026-07-27 through the deployed configuration: the full vision system prompt (Chapter 6 → deployment guide → OWUI configuration), the ragfarm tool schemas attached, tool calls executing for real against mcpo. No harness shortcuts, no missing scaffolding — what runs below is what runs when a board member types the same prompt into OWUI.

Every prompt names the fixture URL up front. All fixtures live on the box at `http://<host>/fixtures/<file>` (served by the `drawio-viewer` `nginx` container on `0.0.0.0:80`) — download and use them directly, or reference them in prompts you author for other tests. The exact prompt string in each example is copy-pasteable.

The current model is `Qwen3-VL-8B-Thinking Q4_K_M` running through `llama.cpp/Vulkan` on the `Radeon 890M` iGPU. Decode is memory-bandwidth-bound at 8 tokens/second on this hardware; every hard number in this chapter is the iGPU floor and scales 6-9× on the dGPU targets in Chapter 4.

1 · RAG over internal corpus — one tool call, full-column table

answer **Fixture:** N/A (RAG queries operate against Qdrant, not against an image) **Model called:** `tool_search_corpus_post` (exactly once — see RULE 1 in the vision system prompt) **Wall-clock:** 540 s, 3473 output tokens **Prompt:**

Dej mi kontakty na projektové vedení v EPC.

Answer (verbatim):

Jméno	Firma	Role/oblast	Tel	E-mail
---	---	---	---	---
Viktor Vážný	EPC	IT, x86, Windows	604 795 022	vit.vazny@epcommodities.cz
Petr Pyzsko	EPC	Sítě	737924492	petr.pyzsko@epcommodities.cz
Michal Sterba	EPC	Aplikace		michal.sterba@epcommodities.cz
Miroslav Laboj	EPC	Aplikace a sítě		miroslav.laboj@epcommodities.ch
Viktor Bobro	EPC	Aplikace		viktor.bobro@epcommodities.cz
Marek Česal	EPC	PM - řízení cut-over procesu	739223474	marek.cesal@epcommodities.cz

Source: `EPC-CutOver_contacts-V1_1.xlsx`

Six contacts, every column preserved, empty phone-number cells left empty (not invented), source .xlsx cited on the last line. **One tool call — the single-tool-per-turn RULE in the vision system prompt now prevents the Thinking model from re-invoking the same tool to "refine".** This is the same category as the Chapter 2 text-preset screenshot for contacts; the vision preset produces the same shape of answer, just slower.

2 · Image classification — "what's on this picture?"

Fixture: <http://127.0.0.1/fixtures/pytorch-dog1.jpg> (97 KB, from pytorch/vision reference set) **Wall-clock:** 38.8 s, 253 output tokens — **fastest turn of the chapter Prompt:**

What is on this picture? Give a one-sentence description of the main subject.

Answer (verbatim):

A Cardigan Welsh Corgi dog sits on a gravel path surrounded by tall pine trees in a forest.

Correct breed identification (Cardigan Welsh Corgi, not the more common Pembroke — the model picked up on the tail), correct posture (sitting), correct surface (gravel), correct surroundings (pine forest). Under 40 seconds — the "wow" turn to open with because it feels instantaneous relative to how the audience is expecting AI-vision demos to lag.

3 · OCR of a hand-written German book page

Fixture: http://127.0.0.1/fixtures/Copilot_20260727_041741.png (1.1 MB) **Wall-clock:** 566 s, 4000 output tokens **Prompt:**

This is a hand-written German book page. OCR every visible line into German verbatim – keep line breaks and spelling as printed.

Answer (verbatim, first 870 characters):

Kapitel 4 – Unterwegs in der Stadt

Lesetext

Maria wohnt seit drei Jahren in Leipzig. Jednn Morgen fährt sie mit der Straßebahn zur Arbeit. Die Fahrt dauert onghau zwonnig Minuten. Auf dem Weg liest sie oft eih Buch oder hört Musik.

In der Mittagspause geht Maria gern spazieren. Sie kauft manklans frisches Obst auf den Markt oder trinkt ejen Kaffee im Park. Am Nachnuung abortet sie an versefrenbent Projeten and devefrentet E-Mails.

Nach Feierabend trifft sie Freunde oder besucht einen Sprachkurs. Um zehn Uhr kommt sie nach Hause und bereitet das Abesnkeer zu.

Wortschatz – die Straßebahn – tram – der Markt – market – beantworten – to reply – die Mittagspause – lunch break – der Markt – market – Feierabend – end of workday.

Übung 1 – Fragen zum Text ...

Structure recovered end-to-end: chapter title, "Lesetext" section, main narrative paragraph, second paragraph, third paragraph, "Wortschatz" (vocabulary) list, "Übung 1" (Exercise 1) heading. Some hand-writing artefacts in the letters (Jednn for Jeden, Straßebahn for Straßenbahn, zwonnig for zwanzig, abortet for arbeitet, Abesnkeer for Abendessen) — **honest OCR imperfection at 8 B parameters on cursive-adjacent hand-writing**. A dGPU-hosted 32 B tightens this significantly; the shape of the answer is already correct. This is not a failure mode; it's the difference between "computer transcribed 95 %" and "computer transcribed 100 %".

4 · draw.io round-trip — reverse every arrow direction

Input fixture: http://127.0.0.1/fixtures/dependency_input.drawio (three nodes: Frontend → API → Database, both arrows left-to-right) **Output fixture:** http://127.0.0.1/fixtures/dependency_reversed.drawio (saved from the model's response, downloadable, opens directly in the local Draw.io editor at <http://127.0.0.1/drawio-editor.html>) **Wall-clock:** 361 s, 2686 output tokens **Prompt:**

The following text between <drawio> tags is a draw.io XML file (also downloadable at http://127.0.0.1/fixtures/dependency_input.drawio). Reverse every arrow's direction and return the modified draw.io XML inside a fenced ```xml block. Do not change any styling, node positions, or IDs – only swap `source` and `target` on each edge cell.

```
<drawio>
<mxfile host="app.diagrams.net">
  ...
  <mxCell id="e1" ... source="A" target="B">
  <mxCell id="e2" ... source="B" target="C">
</mxfile>
</drawio>
```

Answer: a fenced ```xml block containing the exact same document with **only** the two edge cells' source and target attributes swapped:

- e1: **was** source="A" target="B", **now** source="B" target="A"
- e2: **was** source="B" target="C", **now** source="C" target="B"

Every other line (node positions, colors, geometry, IDs, XML boilerplate) preserved byte-for-byte. The diff between input and output is 4 lines — exactly the 2 edges. Verified by opening `dependency_reversed.drawio` in the local Draw.io editor at <http://127.0.0.1/drawio-editor.html>.

This is the pattern. Any structured-file transformation — YAML, JSON, XML, drawio, mermaid — can be round-tripped by describing the transformation in natural language. No custom parser, no ETL job, no engineering work per transformation.

5 · Mermaid generation for a stack diagram (natural-language spec)

Fixture: N/A (pure text-to-mermaid — no image input) **Wall-clock:** 351 s, 2669 output tokens **Prompt:**

Draw a mermaid flowchart of this stack: Frontend calls vLLM; Frontend also calls MCP;
MCP calls RAG; RAG calls Embedder and Reranker in parallel; both feed back into RAG
which returns to MCP. Use mermaid syntax, no draw.io.

Answer (verbatim):

```
graph LR
    F[Frontend] --> vLLM
    F --> MCP
    MCP --> RAG
    RAG --> Embedder
    RAG --> Reranker
    Embedder --> RAG
    Reranker --> RAG
    RAG --> MCP
```

Every relationship in the spec captured: Frontend fans out to vLLM and MCP; MCP goes into RAG; RAG parallels to Embedder + Reranker with return arrows; RAG closes back to MCP. Rendered inline by OWUI as SVG. Same category as the Chapter 2 mermaid tree screenshot — pure text-to-diagram, no image required.

6 • Where the current iGPU hardware hits a wall

Two prompts in this batch hit the harness's 600-second per-request timeout without emitting an answer. The model was generating, just past our client's patience budget. Both would complete on a dGPU that decodes 6-9× faster.

FW-rules RAG query in Czech (Jaká jsou FW pravidla pro host leadb229p.lea.piz?, from Chapter 2 screenshot 1) → timed out. The Chapter-2 screenshot proves this category works on the text preset (Qwen 2.5-7B greedy, sub-90 s); on the vision preset (Qwen 3-VL-8B Thinking, non-greedy, longer reasoning), the same category exceeds the current iGPU's practical live-demo budget.

German-page OCR + full Czech translation (Copilot_20260727_042037.png typeset German → Czech line-by-line) → timed out. OCR alone succeeded on the hand-written page above; adding a full second language as translation doubles the output tokens. On dGPU: estimated 60-90 seconds for this exact prompt.

These are not the model failing; they are the current iGPU's throughput ceiling with a Thinking-class model. Every other example in this chapter also runs 6-9× faster after the dGPU swap — from "usable" to "snappy" to "instantaneous."

What this chapter demonstrates

Five reproducible, board-safe live tests:

- **RAG lookup with structured output** (Section 1) — the RAG-first architecture with single-tool discipline. Same shape of answer as the Chapter 2 text-preset screenshot for contacts.
- **Image classification** (Section 2) — 39-second turn, cleanest wow.
- **Multilingual OCR** (Section 3) — 8B model produces near-perfect German recovery from hand-writing with honest, small character-level artefacts.
- **Structured-file transformation round-trip** (Section 4) — the pattern that generalises to any YAML/JSON/XML editing task. Downloadable output.
- **Text-to-mermaid diagram** (Section 5) — natural-language specification → rendered graph.

Two ceilings honestly named:

- Anything requiring >2000 output tokens through the Thinking reasoning trace at 8 tok/s decode drifts past a 10-minute wall-clock. Fixed by the dGPU swap (Chapter 4).

- Fine-detail OCR (small hand-written cursive, small mono-space code) is where the 8B model shows its parameter count. Fixed by moving up to 32B on the dGPU.

Everything above is a live capability today. Every prompt is copyable from this document into OWUI on the demo box and produces the same shape of answer.

Chapter 4 Hardware requirements (dev + prod ladder)

Hardware requirements — how they scale, and what to buy

The previous chapter demonstrated what an 8-billion-parameter vision-language model does today, on an integrated GPU sharing 30 GB of LPDDR5x with the CPU. The ceiling is memory bandwidth (130 GB/s aggregate → 8 tokens/second decode). Below is the ladder of currently-sold hardware that removes each successive ceiling and what you get for that money.

Two things worth flagging up front:

- **Model sizes have grown dramatically and continue to.** GPT-3 (2020) shipped at 175 B parameters; ChatGPT-original (2022) collapsed the quality bar down to something in the 20-30 B range. Qwen 2.5, Llama 4, Mistral Nemo — 7 B to 30 B was the productive class through 2024. Qwen 3 (2025-2026) and DeepSeek V3 pushed the useful ceiling to 235 B mixture-of-experts. **Every emergent capability jump — reliable tool-calling, coherent multi-turn reasoning, no-hallucination retrieval, per-domain multilingual competence — happens at a specific parameter threshold.** You don't get "a little bit of" tool-calling by half-fitting a model that's a little too big; you get zero or one, depending on whether you crossed the threshold.
- **What "runs" a model changes as you cross size thresholds.** A 7 B fits on any decent workstation. A 30 B needs 24 GB VRAM. A 72 B needs 48 GB. A 235 B MoE with A22B active needs multi-GPU tensor-parallel or a fabric-connected DC-tier card. **The delta from one tier to the next is not linear** — it's an emergence step, and if the customer's workload is on the wrong side of that step, no amount of squeeze on the current tier recovers it.

This is why the hardware discussion below climbs by rung rather than picking a single "recommended spec." Where you land depends on which emergent capabilities your workload requires.

Tier 0 — the original PoC box (HISTORICAL — kept as the bottom rung)

Component	Spec	Notes
Box	AceMagic F5X (€1300)	Mini-PC, 500 mL volume, single 65 W USB-C PD or barrel.

CPU/iGPU/NPU	AMD Ryzen AI 9 HX 370	Strix Point: 12 cores (4× Zen 5 + 8× Zen 5c), Radeon 890M iGPU (gfx1150), XDNA 2 NPU (50 TOPS INT8, currently unused for LLM inference on Linux).
Memory	64 GB LPDDR5x-7500	Shared with iGPU as UMA (BIOS allocates 30 GB nominally, up to 48 GB visible via GTT extension).
Bandwidth	130 GB/s aggregate	The decode ceiling — everything downstream is bandwidth-bound.
Storage	2× 4 TB NVMe Gen 4	Fine for corpus + Qdrant + several models on disk.

What runs here well: Qwen 2.5-7B greedy for text, Qwen 3-VL-8B-Thinking for vision, BGE-M3 embedder on CPU, bge-reranker on iGPU. **8 tokens/s decode, 170 tokens/s prefill.** OCR of a dense receipt takes 40 s; a Farsi document with translation takes 3 minutes.

What doesn't run here: anything above 10 B parameters at Q4 (VRAM budget), multi-user concurrent load (single generation slot bottleneck), 32k context with -np 4 parallelism (RAM squeeze), any 30 B-class model where the emergent-capability step lives.

Cost profile: €1300 box, 35 W idle, 65 W under load. Effectively free to run 24/7.

This box is retired. Everything above is preserved deliberately: it is the bottom rung of the ladder and the only tier where we have lived with the consequences of being under-provisioned. The numbers in the next chapter that quote 8 tok/s and three-minute Farsi OCR all come from here.

The one formula — how to size any of this

Everything below is arithmetic on one relationship, and we have measured the constant rather than borrowed it.

Decode is memory-bandwidth-bound. Generating a token requires reading the weights it activates out of memory. So:

$$\text{tokens/second} \approx \text{efficiency} \times \text{memory_bandwidth} \div \text{bytes_read_per_token}$$

and bytes_read_per_token is a function of **active** parameters, not total. That distinction is the single most important fact on this page.

The constant, back-solved from our own hardware

Two models measured on the DGX Spark, 2026-08:

model	active bytes/to-ken	measured tok/s	⇒ effective band-width
Qwen3-VL-30B-A3B FP8 (MoE, 3 B active)	3 GB	68	204 GB/s
Qwen3-VL-32B FP8 (dense, all 32 B active)	33 GB	5.9	195 GB/s

Two independent models landing within 5% of each other: **200 GB/s effective, about 73% of the Spark's 273 GB/s specification.** Apply that 0.73 to any vendor bandwidth figure and the estimate holds.

The first consequence: never buy dense

Those two models are the same size class and differ by **11.5×** in speed. The MoE reads 3 GB per token; the dense one reads 33 GB. Every euro of a hardware budget should go to bandwidth and to mixture-of-experts models, not to parameter count. A dense 70 B on this class of machine is not slow — it is unusable, at roughly 2-3 tok/s, which is an hour and a half for one complex answer.

Sizing a target

Rearranged, for a model you are considering:

bandwidth needed = active_bytes_per_token × target_tok/s ÷ 0.73
 VRAM needed = weights + KV cache + ~3 GB overhead, per resident model

Worked, for the first genuinely larger planner class — Qwen3-VL-235B-A22B, 22 B active:

quantisation	bytes/token	for 60 tok/s	for 100 tok/s	weights
NVFP4	12 GB	990 GB/s	1.6 TB/s	130 GB
FP8	22 GB	1.8 TB/s	3.0 TB/s	235 GB

Bandwidth buys latency. VRAM buys concurrency. They are separate purchases and a proposal that confuses them will be wrong in one direction or the other.

Tier 0.5 — NVIDIA DGX Spark (the current substrate)

What the project actually runs on today, and what every measurement in this bundle was taken on.

Component	Spec	Notes
Box	NVIDIA DGX Spark (GB10 Grace Blackwell)	Desktop form factor.
Compute	sm_121, FP4 tensor cores	356 TFLOPS FP4 measured on the flashinfer b12x path.
Memory	128 GB unified (121.7 GiB usable)	CPU and GPU share it. There is no separate VRAM to hide in.
Bandwidth	273 GB/s specified, ** 200 GB/s measured effective**	THE constraint. See the formula above.
Networking	ConnectX-7, 200 Gb/s dual QSFP	Two Sparks can be linked.

What runs here well: two 30 B-class MoE models resident simultaneously — 0.70 of a 0.72 memory-budget ceiling — switchable mid-conversation, at **68-76 tok/s** each. The full RAG stack (BGE-M3 embedder, bge-reranker-v2-m3 cross-encoder, Qdrant, seven containers) alongside them. Vision, tool calling, in-chat diagram rendering, code execution.

What does not run here: any dense model above 30 B at a usable speed; 70 B-class anything; a 235 B MoE at any quantisation (NVFP4 alone is 130 GB against 121.7 GiB usable). Serious concurrency — the budget is spent on having two models resident rather than on KV cache for many users.

Why it was still the right first buy: the entire architecture — Open WebUI, mcpo, the MCP toolchain, hybrid retrieval, reranking, guarded actuation — is built and proven here and redeploys onto any tier above unchanged. Only OPENAI_API_BASE_URL moves. A €200 k lab does not teach you which prompt rule breaks tool calling; a year of using this one did.

Tier 1 — Developer workstation

A single-user developer machine — one person builds, tests, and demonstrates on it. Removes the tokens/second ceiling enough to iterate on prompts and tool chains without being interrupted every 40 seconds waiting for a Thinking-class model to finish.

Component	Recommended	Alternatives (currently sold)
GPU	NVIDIA RTX 5090 · 32 GB GDDR7 · 1790 GB/s (€2 400)	RTX PRO 5000 Blackwell (48 GB GDDR7, workstation-flavor, €4 800); AMD Radeon PRO W7900 (48 GB, €3 800; ROCm-only, worse LLM ecosystem).
CPU	AMD Ryzen 9 9950X3D (16-core, 3D V-Cache)	Intel Core Ultra 9 285K (24-core). Either fine — LLM decode is GPU-side; CPU only matters for RAG orchestration + tool calls.
RAM	128 GB DDR5-6400 (2× 64 GB, dual-channel)	96 GB or 192 GB fine. Not the bottleneck.
Storage	2× 4 TB NVMe Gen 5 (one system, one models/corpus)	8 TB total roughly fits a "several models on disk + full corpus + Qdrant" without pruning.
PSU / case	1000 W 80+ Platinum · full-tower	RTX 5090 pulls 575 W peak, wants proper airflow.
Approx BOM	€3 500 - €4 200 (5090 build)	5090 + 9950X3D + 128 GB + 8 TB NVMe + PSU + case + cooling.

What runs here well: Qwen 3-VL-8B-Thinking at **50-70 tokens/s decode** (7× the PoC), the full 32k context comfortably, Qwen 3-VL-30B-A3B (a 30 B mixture-of-experts with 3 B active — fits in 32 GB Q4) at 15-25 tokens/s, live iteration cycles feel instant.

What still doesn't run here: 70 B dense models at any usable quant, multi-user concurrency beyond 3-4 parallel slots, deep-batching for high-throughput scenarios (agentic workflows spawning 20+ parallel tool calls).

Best value point. The 5090 is roughly the last consumer-grade rung before workstation/DC pricing kicks in. If a workload doesn't require 48 GB VRAM (i.e. you're happy with 30 B-class MoE and don't need multi-user), the 5090 is the sweet spot.

Tier 2 — Small-team production server (10-50 concurrent users)

The point at which a workstation stops being enough. One or two workstation-class GPUs in a rackmount chassis, driving a departmental or small-customer deployment. Serves both the LLM and the reranker + embedder from the same server; Qdrant colocated.

Component	Recommended	Alternatives
GPU (×1 or ×2)	NVIDIA RTX PRO 6000 Blackwell Max-Q · 96 GB GDDR7 · 1800 GB/s (€10 500 each)	NVIDIA L40S (48 GB GDDR6, €8 500 — older Ada arch, less memory but well-established) · AMD MI300X (192 GB HBM3, €15 000 — much cheaper per GB VRAM but ROCm software friction remains a real cost).
Server chassis	Supermicro AS-2124GQ-NART or Dell PowerEdge R760xa	2U or 4U depending on GPU count; needs PCIe 5.0 x16 per card and 800 W+ per card.
CPU	1× AMD EPYC 9354P (32-core, PCIe 5.0 lanes)	Or Intel Xeon 6 Granite Rapids equivalent. Single-socket is enough.
RAM	512 GB DDR5-4800 ECC RDIMM	For Qdrant + FS cache + inference server + head-room.
Storage	4× 8 TB NVMe Gen 5 U.2/ E1.S (RAID 10)	16 TB usable — corpus + Qdrant snapshots + N models on disk.
Networking	2× 25 GbE (or 100 GbE if fronting a bigger cluster)	For OpenNebula integration + client-facing traffic.

PSU	2× 2000 W redundant	RTX PRO 6000 = 600 W each; leave margin.
Approx BOM	€18 000 - €25 000 (single-GPU) · €28 000 - €35 000 (dual-GPU)	Chassis + CPU + RAM + storage + GPU(s) + PSUs.

What runs here well: Qwen 3-72B-Instruct at Q4 (fits in 96 GB with room), Qwen 3-VL-32B-Thinking at BF16 (fits in one 96 GB card), 8-16 parallel inference slots, sub-second first-token latency, real customer-facing SLAs. Full ADR-0008 reranker at BF16 not Q4.

What still doesn't run here: Frontier-class 405 B+ dense models, or 671 B MoE like DeepSeek V3, or multi-tenant workloads where a hundred concurrent users each expect an interactive response.

Best fit: internal-tool deployments (single company, few dozen active users), customer-facing pilots, embedded-in-product AI features for SMEs.

Tier 3 — DC-tier serious production (100s of concurrent users, or a

foundation model to serve) At this rung we're in HGX/DGX territory — GPU-fabric-connected servers with NVLink, deployed in ones or twos.

Component	Recommended	Alternatives
GPUs	NVIDIA HGX H200 · 8× H141 · 1128 GB HBM3e total · 4.8 TB/s per card · NVLink 5 (€300 000 for a bare HGX baseboard alone)	AMD MI325X 8-way system (€250 000, comparable HBM3e VRAM, PCIe fabric not NVLink); NVIDIA B200 HGX (Blackwell DC — €400 000, another generation newer).
Server platform	Supermicro SYS-821GE-TNHR or NVIDIA DGX H200	8U to 10U form factor; requires 400 V 3-phase power, liquid cooling, dedicated rack.
CPU	2× Intel Xeon Platinum 8592+ (or dual AMD EPYC 9754)	128 cores total; provides PCIe lanes + RAM channels for orchestration.

RAM	2 TB DDR5-4800 ECC (24× 96 GB)	Feeds the CPU-side vLLM prefetch buffers + KV-cache spillover.
Storage	60 TB NVMe Gen 5 (RAID 10) + separate parallel filesystem for the model catalogue	For continuous training-data ingestion in addition to inference.
Networking	2× 400 GbE ConnectX-7 (or InfiniBand NDR)	For multi-node scaling and client-facing throughput.
Power	10 kW at load	Requires proper DC facility.
Approx BOM	€400 000 - €600 000 for a single HGX/DGX server	Ready-to-run; excludes rack, PDU, cooling.

What runs here well: Qwen 3 Coder 480B-A35B-Instruct at BF16 (needs multi-GPU tensor-parallel), Qwen 3-VL-235B-A22B, DeepSeek V3 671B, Llama 4 Maverick 400B — the frontier open-weight tier. Hundreds of concurrent users. Deep-batching. Fine-tuning by LoRA/QLoRA. Continuous embedding backfill.

What doesn't run here: frontier proprietary models where the weights aren't yours to load. Otherwise there is no ceiling in the currently-shipping open-weight space.

Best fit: the point at which the product is the AI, not "a product with AI in it." Enterprise SaaS pivots, foundation-model resellers, sovereign-AI national initiatives.

The trajectory that matters for planning

The parameter-count / VRAM / bandwidth requirements have grown roughly **2×-3× per year** across every open-weight release wave since 2022, with no plateau in sight. The Qwen 2.5 → Qwen 3 → Qwen 3-VL step alone doubled the useful-model class from 7 B to 32 B; the next step will likely reach 72 B-dense or 200 B-MoE as the "productive floor." Hardware planned today for a 3-year lifecycle should be provisioned with **2× headroom** on VRAM and bandwidth, not sized to today's model.

Consequences for this project's roadmap:

- **Tier 0 (PoC)** is a genuine on-prem demonstration platform and will remain useful for prompt-engineering, corpus ingestion tuning, and travel demos. It will not track the model frontier.
- **Tier 1 (single dGPU workstation)** is what unlocks Qwen 3-VL Instruct + Thinking at interactive latencies + the 30 B-A3B MoE class. This is the minimum plausible dev environment going forward. Any customer-facing demo run on Tier 0 hardware will visibly lag against expectations set by cloud AI.
- **Tier 2 (workstation-class rack GPU)** is the first customer-deployable tier — the point at which a company can host their own instance without knowing what they're doing at the infrastructure level. The whole ragfarm architecture (Open WebUI + mcpo + MCP toolchain + hybrid retrieval + rerank + guarded actuation) is designed to redeploy unchanged onto this tier; only the `OPENAI_API_BASE_URL` moves.
- **Tier 3 (DC HGX)** is the tier where a hosted-AI product becomes competitive with commercial offerings on both quality and speed. Not required for the current use case; noted for the roadmap.

The single most important decision to defer or make: whether Tier 1 (single-workstation dGPU) makes sense as an interim step before jumping to Tier 2 for customer deployment. On this exact project, Tier 1 unblocks development, and Tier 2 unblocks the customer pilot. The €4 000 vs €25 000 delta is small compared to the year of engineering time each unlocks.

The concrete ask, if you have one conversation to make it in

Two asks come up in practice: make the current models three times faster, and run a planner that is not 30 B-class. They cost very differently.

3× on the current MoE (68 → 200 tok/s) needs $3 \text{ GB} \times 200 \div 0.73 \approx \mathbf{820 \text{ GB/s}}$. That is one workstation card. It is the cheap ask.

A real planner is what sets the bill, because the 235 B-A22B class reads seven times more per token.

ask	hardware	what it buys
Development un-block	1× RTX PRO 6000 Blackwell · 96 GB · 1.8 TB/s · €10-12 k	3× on the current models, a 120 B-class MoE. Not the 235 B.

The planner	2× RTX PRO 6000 · 192 GB · €20-25 k GPU, €35-40 k built	235 B-A22B at NVFP4 (130 GB) with room for KV, near 110 tok/s — past the 3× ask on a model 7× the active size.
Customer pilot	2× H200 · 282 GB · 4.8 TB/s · NVLink · €70-100 k	The same planner at full FP8, 160 tok/s, and enough KV for real concurrency.

Three honesty flags for whoever quotes this:

- **Only the 200 GB/s is ours.** Every other bandwidth figure is a specification. Apply the measured 0.73 efficiency and the estimates hold, but they are estimates.
- **Prices are indicative, not quotes.** The ratios matter more than the absolutes; get real numbers before committing.
- **Multi-GPU is not free.** A PRO 6000 pair talks over PCIe, not NVLink, so tensor-parallel on a 130 GB model gives some of that 110 tok/s back. H200s do not have that problem, and that is the argument for the more expensive tier if the planner is the point.

Chapter 5 Deployment guide

Deployment & prod re-deployment notes

Author: David Kubicek (david.kubicek@eywo.cz)

Operational facts needed to run and redeploy the system. **ADR-0013/0003 hold the why this holds the concrete what.**

Hardware note. This system now runs on an **NVIDIA DGX Spark (GB10, sm_121, 128 GB unified memory)**. It was prototyped on an AMD Ryzen MiniPC (iGPU + Vulkan + GGUF), and that migration is complete — ADR-0013 supersedes ADR-0001. AMD-specific mechanics have been removed from this document; what is deliberately kept is (a) **model-behaviour knowledge**, which is hardware-independent and was expensive to learn, and (b) the **AMD baseline performance table**, retained as the before/after comparison. Both are marked where they appear.

Runtime topology — live component registry

The maintained inventory of everything that runs. This is the **live runtime registry** it supersedes the point-in-time snapshot in ADR-0005 as "what is running now." The naming, port-allocation and collocation **rules** (how a short name derives its directory / container / unit / port / mount) stay authoritative in **ADR-0005** — this table is the current state, that ADR is the policy. Add a component → add its row here and take the next 81xx port per ADR-0005. All container services run `network_mode: host` (see below) and bind loopback except open-webui.

plane	component	directory	compose service / container	bind	mcpo mount	agent-facing surface	mutates
host	vllm	.venv-vllm	— (sys127.0.0.1:8080 temd ragfarm-vllm)	127.0.0.1:8080	—	OpenAI base URL (swap- pable)	no
host	em-bedder	services/ embedder	— (sys27.0.0.1:8090 temd	27.0.0.1:8090 / embed	— nal —	inter-dense+sparse	no

			ragfarm-embedder)			embed, CUDA	
host	reranker	(built at ~/llama.cpp)	— (systemd ragfarm-reranker)	127.0.0.1:8081 / reranking	—	internal — bge-reranker-v2-m3 GGUF, CUDA (ADR-0008)	no
host	ingester (+ watcher)	services/ingester	— (systemd ragfarm-ingester-watcher)	—	—	batch + autonomous incremental sync (ADR-0006)	writes Qdrant
container	qdrant	up-stream image	qdrant infra-qdrant	127.0.0.1:6333 / 6334	—	retrieval store; volume qdrant_data	no
container	rag	services/rag-rag-retrieval	rag-retrieval infra-rag-retrieval	127.0.0.1:8104 / rag	—	tool server — search_corpus	no
container	placement	services/mcp-placement	mcp-placement infra-mcp-placement	127.0.0.1:8101 / placement	—	tool server — mock until Open-Nebula	no
container	host-control	services/mcp-host-control	mcp-host-control / infra-mcp-host-control	127.0.0.1:8102 / host-control	—	wrapper only — reboot_guarded (ADR-0004)	YES
container	fs	services/mcp-fs	mcp-fs infra-mcp-fs	127.0.0.1:8103 / (unbridged)	—	experimental — not in the	no

						agent path	
con- tainer	mcpo	up- stream image	mcpo infra- mcpo	127.0.0.1:8000	(is the front)	the single Ope- nAPI tool end- point	n/a
con- tainer	open- webui	up- stream image	open- webui / infra- open- webui	0.0.0.0:3000	—	the UI (only LAN- ex- posed; auth- gated)	n/a

ragfarm-stack.service launches the container plane **except** mcp-fs (unbridged) and **except the ingester** (a host job / the watcher unit, not the stack): qdrant, rag-retrieval, mcp-placement, mcp-host-control, mcpo, open-webui.

Three processes share the one GPU and the one memory pipe — there is no separate VRAM on GB10 (ADR-0013 §5). Measured while idle-but-loaded:

process	GPU memory
vLLM (:8080, LLM_GPU_MEM_UTIL=0.50)	59.5 GB
embedder BGE-M3 (:8090)	1.6 GB
reranker llama.cpp (:8081)	0.9 GB
total	62 GB of 121 GB

llama-server now runs **once**, for the reranker only (:8081 --reranking, ADR-0008/0013 §3) — generation moved to vLLM. The embedder is embeddings-only. Argue any "let's also run X" against the table above, not against a spare-capacity assumption.

Why host networking (load-bearing PoC fact)

The LLM (:8080) and embedder (:8090) are host processes bound to **127.0.0.1 only** (not exposed to the LAN). A bridge-network container cannot reach a loopback-bound host service via host.docker.internal (connection refused). So

rag-retrieval / mcpo / open-webui run with `network_mode: host` and reach the host services at 127.0.0.1. They keep their own binds on loopback (`RAG_HOST=127.0.0.1`, `mcpo --host 127.0.0.1`) except open-webui, which binds 0.0.0.0 for remote access.

Remote UI access

`OWUI_HOST` (compose env) controls the Open WebUI bind: 0.0.0.0 (default now = remote access, login-gated by `WEBUI_AUTH`) or 127.0.0.1 (loopback-only; reach via `ssh -L 3000:127.0.0.1:3000 <host>`). **Restrict :3000 at the host firewall to trusted networks** — it's the only externally reachable service.

Tested model versions (known-good baseline)

The fetch scripts default to **latest** (easy swapping/experimentation), but the model set every gate + eval in this repo was validated against is the pinned one below. Pass the `--revision` flags to reproduce it exactly; omit them for latest. These pins are for **reproducibility only** — not a security constraint (the old no-pickle rule is retired, see the model-format note in `scripts/lib-models.sh`).

role	model	deployed revision	reproduce with
LLM	ig1/Qwen3-VL-72B-A3B-Instruct-NVFP4 (NVFP4 safetensors, 17.9 GiB, vision)	13d26f008628eebe9b435559b4302	Step 302 fragment in <code>scripts/deploy.sh</code> (<code>snapshot_download</code>)
embedder	BAAI/bge-5617a9f61	b028005a4858fdac845db406aefb181 (latest; this repo ships only <code>pytorch_model.bin</code> — it has no <code>model.safetensors</code>)	<code>scripts/fetch-encoder.sh</code>
reranker	BAAI/bge-reranker-v2-m3	latest, converted to f16 GGUF locally	<code>scripts/fetch-encoder.sh</code>

Served LLM alias is `qwen3-vl-30b-a3b` and it is **load-bearing**: `infra/openwebui/setup_openwebui.py` keys `MODEL_TUNING` by the alias advertised on `/v1/models`, so serving under another name makes the `ragfarm-vision` preset silently fail to bind.

The embedder + reranker must stay a **compatible pair** (`fetch-encoder.sh --list`). Per-model detail lives in `models/{llm,embeddings,reranker}/MODEL.md` — including the measured NVFP4 MoE backend findings, which contradict parts of ADR-0013 and are worth reading before touching the serving flags.

Older baseline, for reference only: the PoC ran Qwen2.5-7B-Instruct Q4_K_M GGUF with bge-m3 at 50f9396f.... Neither is deployed now; GGUF is the retired llama.cpp generation path.

Autostart & lifecycle

What starts on boot

- **Host services** (systemd, already enabled): ragfarm-vllm.service (LLM, CUDA), ragfarm-reranker.service (CUDA cross-encoder, ADR-0008; Nice=5), ragfarm-embedder.service (CUDA embeddings), ragfarm-ingester-watcher.service (autonomous corpus sync, ADR-0006).
- **Container stack:** ragfarm-stack.service runs `docker compose up -d` for the six container services (see the topology note above), then its `ExecStartPost` runs `scripts/mcpo-heal.sh`. Containers also carry `restart: unless-stopped` and `docker.service` is enabled — so a normal reboot restarts them anyway; the unit guarantees ordering + one control point (install steps in the unit header).

Start / stop / status everything (host, as dave)

One command (recommended) — `scripts/stack.sh` runs the whole ordered sequence and health-checks the result. Reach for this unless you're debugging one service:

```
scripts/stack.sh start      # host services → container stack → health check
scripts/stack.sh stop       # container stack → host services (reverse order)
scripts/stack.sh restart    # stop, then start
scripts/stack.sh status     # systemd unit + container status at a glance
scripts/stack.sh health     # probe every endpoint; non-zero exit if any is down
```

The explicit per-service sequence below does exactly the same thing — use it when you need to bring one piece up or down on its own.

Cold start — host model hosts first, then the container stack:

```
sudo systemctl start ragfarm-vllm ragfarm-reranker ragfarm-embedder ragfarm-
ingester-watcher
sudo systemctl start ragfarm-stack
```

Stop everything — stack first, then host services:

```
sudo systemctl stop ragfarm-stack
sudo systemctl stop ragfarm-ingester-watcher ragfarm-embedder ragfarm-reranker
ragfarm-vllm
```

Status at a glance:

```
systemctl --no-pager status ragfarm-vllm ragfarm-reranker ragfarm-embedder
ragfarm-ingester-watcher ragfarm-stack
docker compose -f infra/compose.yaml ps
```

Restarting individual pieces (the gotchas)

- **rag-retrieval**: any restart severs mcpo's streamable-http MCP session — always follow with `scripts/mcpo-heal.sh` (or just `sudo systemctl restart ragfarm-stack`), otherwise tools come up unmounted (ADR-0007 note #4; the boot healer is the stack unit's `ExecStartPost`).
- **reranker & embedder**: independent host services (a CUDA `llama-server` on `:8081` and the CUDA embedder on `:8090`); restarting one doesn't touch the other. The reranker loads its 1.1 GB GGUF in 1-2 s; rag-retrieval reaches it via `RERANK_ENDPOINT (:8081/reranking)`. One gotcha: `llama.cpp` reranking scores each (query,doc) pair in one physical batch, so the unit sets `-b/-ub 4096` — if chunks ever grow past that, raise it (see `ragfarm-reranker.service`).
- **vLLM cold start is slow, and that is not a hang**. First start after a cache miss JIT-compiles FlashInfer kernels: measured **813 s cold vs 3 min warm**. Two things make this survivable and both live in `.env`: `MAX_JOBS` (caps parallel CUDA compiles — unset it defaults to `nproc+2` and the OOM killer takes the service at 98 GB) and `LLM_GPU_MEM_UTIL` (steady-state weights+KV budget). See `.env.example` and ADR-0013 §5a — they are different budgets, and lowering the vLLM memory flags does nothing for the cold-start peak. The unit has `StartLimitBurst=3`, so a genuinely broken config goes failed instead of crash-looping while `systemctl is-active SAYS` activating.

GPU contention & the reranker's priority (tunable — revisit with

real load) `ragfarm-reranker.service` runs at `Nice=5`, a deliberately mild de-prioritisation so the interactive LLM wins contention (ADR-0013 §3). It is set low on purpose: we are **single-user today**, and the point is to observe how much these three actually fight over the GPU before tuning harder.

Know what Nice does and does not buy. It is CPU scheduling only — it does not yield GPU time. Kernel ordering on the device is the driver's business, and nothing in the unit deprioritises the reranker's CUDA work. What it does help with is the host-side half: HTTP handling, tokenisation, and the batch marshalling around each rerank call.

To tune (any of these, no code changes):

- raise/lower `Nice=` in `manifests/ragfarm-reranker.service`, reinstall, restart;
- if GPU contention is the real problem, the levers are CUDA MPS priorities, or simply not running a bulk re-rank concurrently with interactive generation;
- if the reranker is starving the LLM of memory rather than compute, lower `LLM_GPU_MEM_UTIL` and re-check the topology table above.

Measure before tuning: a rerank turn was 1.9 s on the old iGPU and the model is unchanged, so if it regresses noticeably under concurrent load, that is the signal — not a number anyone should guess at up front.

- Manual docker compose ops need the proxy env first: `source scripts/proxy-env.sh` (image pulls + container proxy inheritance).

Operator scripts (scripts/) — a newcomer's map

Everything here runs from the repo root on the host as dave. Grouped by job.

Deploy & lifecycle

- `stack.sh` — the **one-command operator entry point**: `stack.sh {start|stop|restart|status|health}`. Brings the whole stack (host llama/reranker/embedder/ingester-watcher units + the container stack) up or down in the right order and health-checks every endpoint. Use this for day-to-day start/stop; the per-service systemd sequence is only for debugging one piece (see Autostart & lifecycle above).
- `deploy.sh` — the reproducible, idempotent full deploy of the durable stack: ordered phases (preflight → venv → host services → stack → corpus → watcher → verify), each ending in a machine-checkable gate; safe to re-run. Uses `sudo` only for the specific systemd install/enable actions (never wraps the whole script). Picks a Python dependency profile — `--profile cpu` (default) or `cu13` — pinned in `services/requirements.lock / requirements.cu13.lock` (identical package set; only the torch wheel index differs, CPU vs CUDA 13.x). Builds the reranker GGUF if absent and installs all host units incl. `ragfarm-reranker`. The AI-out-of-the-loop path for repeatable deploys.
- `mcpo-heal.sh` — waits until the MCP backends accept TCP, then restarts mcpo once so every tool mounts cleanly (works around the streamable-http boot race). Runs as the stack unit's `ExecStartPost`; run it by hand after any ad-hoc `rag-retrieval` restart.
- `proxy-env.sh` — **source it, don't execute**, before any network command (`pip / HuggingFace / docker compose`); loads repo-root `.env` and normalizes `HTTP(S)_PROXY/NO_PROXY` (guarantees loopback + containers bypass the proxy).

Model management (fetch / hot-swap)

REPLACED 2026-08-05 by `scripts/fetch_llm.py` + `scripts/activate_llm.py`, driven by the git-tracked registry `models/llm/active.json` (ADR-0013 §2a).

```
scripts/fetch_llm.py -m Firworks/Qwen3-VL-32B-Thinking-nvfp4 # register + download
scripts/fetch_llm.py --sync # fetch everything the registry lists as missing
scripts/fetch_llm.py --verify # byte-check what is on disk against the Hub
```



```
scripts/activate_llm.py -s 0 -a qwen3-vl-30b-a3b-thinking-nvfp4 # bind slot 0
scripts/activate_llm.py --status # slots, ports, GPU budget
```

SLOTS. vLLM serves ONE base model per process, so two resident models means two processes: ragfarm-vllm@N.service on port 8080 + 2N (8081 is the reranker), each with its own .env.slotN written by activate_llm.py. This is what allows switching model **mid-chat** in Open WebUI with the conversation intact. --gpu-memory-utilization is per process and instances do NOT coordinate — two slots at the single-model 0.50 will OOM; activate_llm.py derives each slot's share and refuses to exceed a 0.72 total.

activate_llm.py errors out on a registry with duplicate model/alias/preset rather than auto-fixing it: those names show up in the UI.

The GGUF-era text below is kept because the mechanics are still accurate for llama.cpp, which still serves the reranker. fetch-encoder.sh is NOT retired.

fetch-llm.sh and activate-llm.sh are RETIRED for generation (ADR-0013). They are GGUF/--mmproj/llama.cpp-shaped, the Spark serves NVFP4 safetensors on vLLM, and deploy.sh no longer calls them. Kept below because the mechanics are still accurate for llama.cpp and no vLLM-shaped replacement exists yet. To change the served LLM today: edit LLM_MODEL_PATH/LLM_SERVED_NAME in .env and restart ragfarm-vllm. fetch-encoder.sh is NOT retired — it still fetches the embedder+reranker pair and is called by deploy.sh.

- fetch-llm.sh — fetch/swap the generative LLM GGUF into models/llm/<slug>/ and write LLM_GGUF_PATH to .env (the unit reads it). --list shows known-good tool-calling LLMs; --repo/--file swap; --force re-fetch. Restart ragfarm-llama. **Vision models** (e.g. Qwen2.5-VL) are handled automatically, no separate flag: a vision-language GGUF needs a second, small "multimodal projector" GGUF passed to llama-server as --mmproj, or it loads but can't see images. Before downloading, the script checks whether the HF repo itself hosts a *mmproj*.gguf; if so it fetches that too and writes LLM_GGUF_MMPROJ, else (plain text model) it **clears** that var — so switching back to a text-only model can't launch with a leftover --mmproj from a previous vision model. Example (fetches both files):

```
scripts/fetch-llm.sh --repo ggml-org/Qwen2.5-VL-7B-Instruct-GGUF \
  --file 'Qwen2.5-VL-7B-Instruct-Q4_K_M.gguf'
```

- activate-llm.sh — switch the active LLM among models **already on disk**, no re-download. Lists every models/llm/<slug>/ that has a complete GGUF (interactive numbered prompt, or --dir <slug> / --path <file> non-interactively; --list prints and exits), and writes LLM_GGUF_PATH + LLM_GGUF_MMPROJ to .env using the same mmproj auto-detect/clear rule as fetch-llm.sh. This is the tool for

"I already fetched three models, which one is live" — fetch once per model, activate freely between them:

```
scripts/activate-llm.sh --list           # what's on disk
scripts/activate-llm.sh --dir qwen2.5-32b-instruct-gguf # switch to it
sudo systemctl restart ragfarm-llama    # RETIRED: unit no longer exists
```

- `fetch-encoder.sh` — fetch/swap the **matched** embedder+reranker pair into `models/embeddings/<slug>/` and `models/reranker/<slug>/` (converts the reranker to GGUF), writing `EMBED_MODEL_PATH + RERANK_MODEL_PATH` to `.env`. `--list` shows compatible pairs. Restart `ragfarm-embedder` `ragfarm-reranker`. Both fetch latest, auto-pick the fastest weight format, and are the ONE place download logic lives — `deploy.sh`'s `venv` phase calls them (no duplication). `lib-models.sh` is their shared helper (sourced, not run).

Activating a swapped model — the fetch/activate scripts only write `.env`; the units read it on (re)start. What you must do after a swap depends on which model changed:

- **LLM** → edit `.env` (`LLM_MODEL_PATH`), then `sudo systemctl restart ragfarm-vllm`. Nothing else — the LLM doesn't touch embeddings or the corpus. (If the model alias changed, also re-run `infra/openwebui/setup_openwebui.py` so the OWUI preset points at There is no `--mmproj` step any more: the old llama.cpp unit attached a separate multimodal projector GGUF, whereas the vLLM checkpoint carries vision natively. `LLM_MMPROJ_PATH` / `LLM_GGUF_MMPROJ` are inert on the Spark.
- **Reranker only** → `sudo systemctl restart ragfarm-reranker`. Query-time only; no re-ingest.
- **Embedder** (`fetch-encoder.sh`) → the vector space changes, so you **MUST** re-embed: `sudo systemctl restart ragfarm-embedder ragfarm-reranker` **then** `.venv/bin/python services/ingester/ingester.py --recreate --corpus /data/corpus` (or `scripts/deploy.sh --recreate-corpus`). **Skipping the re-ingest silently breaks retrieval** — old stored vectors vs new query vectors. `scripts/stack.sh restart` restarts every service but does **not** re-ingest; that step is always manual. (`fetch-encoder.sh` prints this reminder when it actually re-fetched the embedder.)

Retrieval / RAG debugging

- `rag_pool_inspect.py` — dump the first-stage RRF candidate pool for a query (`--branches` adds the dense-only and sparse-only lists). Separates a RANKING problem (chunk is in the pool but scored low → the reranker's job) from a

RECALL problem (chunk never entered the pool → needs better first-stage recall / query expansion). This is the raw, pre-rerank view.

```
.venv/bin/python scripts/rag_pool_inspect.py --branches "hsmbvxi001ts" "proj  
vedoucí EPC"
```

- ingest-embed-test.sh — quick smoke test: Qdrant points carry non-empty sparse vectors, and a known-hostname search_corpus lookup returns rows.

Agent / tool-routing debugging

- dump_mcp_openapi.py — enumerate every mcpo mount and dump its openapi.json plus the exact operationId the model is shown (e.g. tool_search_corpus_post). Ground truth when writing routing hints or debugging why the model does/doesn't call a tool. --full for the whole schema per mount.
- trace_tool_calls.py — drive llama-server with the **real** OWUI tool schemas (pulled live from mcpo) and the deployed grounding prompt at deterministic settings, feeding canned tool results, to watch which tools the 7B calls with what arguments across rounds. Regression-check routing after a prompt/schema change.
- agent.py — headless multi-turn agent over the same stack (llama-server + mcpo tools + the deployed grounding prompt) but with a context loop **we own and measure**. Unlike trace_tool_calls.py it EXECUTES tools for real (incl. a CLI reboot confirm mirroring reboot_guarded.py) and carries real history. Modes: interactive REPL, one-shot prompts, or scripted benchmarks (--scenario reboot-canary). Per turn it splits and times the **deliberate / tool / answer** stages (prefill vs decode each) and reports context size + how many old tool results were elided. This is the control for isolating an OWUI-loop bug from a model/stack bug — e.g. it reproduced the repeated-reboot miss with no OWUI in the loop, proving that failure is model discipline, not compaction.

Open WebUI configuration (reproducible, not hand-clicked)

OWUI stores config in its openwebui_data volume. Recreate it with infra/openwebui/setup_openwebui.py (idempotent) — this pushes TWO presets in one pass (text + vision, see ADR-0009 and the "Vision engine" section below):

- **Tool servers** TOOL_SERVER_CONNECTIONS → /rag (server:0)
 1. /placement (server:1), both path=openapi.json,

auth_type=none.
- **Preset ragfarm** (text): base qwen2.5-7b-instruct + rag/placement/reboot_guarded attached (meta.toolIds) + params.function_calling=native + greedy sampler (temp=0, top_k=1, seed=42) + the RULE-1.5 grounding prompt.

Loadbearing: without it the 7B answers generic prose even when the exact chunk was retrieved.

- **Preset ragfarm-vision (VL):** base auto-detected via `/v1/models` (first entry with capability `multimodal`, overridable via `VISION_BASE_MODEL_ID`), non-greedy sampler (`temp=0.6`, `top_k/top_p/min_p/seed` DROPPED so llama.cpp defaults apply — required by Qwen3-VL Thinking) + vision + file-upload capabilities on
1. a RULE-1..6 prompt that adds image-input rules and a draw.io HTML

template pointing at the local viewer (`127.0.0.1:8091`, see `drawio-viewer` service in `infra/compose.yaml`).

- Capabilities matrix (both presets): `file_context`, `file_upload`, `web_search`, `code_interpreter`, `citations`, `status_updates`, `usage`, `builtin_tools` all ON; `image_generation`, `terminal` OFF; vision only on VL preset.
- Default features (per-chat pre-selected): `web_search` + `code_interpreter`.
- Builtin tools: everything ON except knowledge and calendar (OWUI opt-out convention — absence = enabled).
- OpenAI endpoint: `OPENAI_API_BASE_URL=http://127.0.0.1:8080/v1` (compose env).
- Open WebUI's built-in document RAG is deliberately unused for the corpus (Option B).

Run it (admin token, or email + password on the CLI):

```
OWUI_URL=http://127.0.0.1:3000 OWUI_TOKEN=<admin JWT> \  
.venv/bin/python infra/openwebui/setup_openwebui.py
```

Only one llama-server model is loaded at a time — the preset whose `base_model_id` matches the wrapper's `--alias` is the one that actually works right now; the other stays as stored config waiting for the next model swap.

Verifying the toolchain

`infra/openwebui/check_toolchain.py` — exit 0 full pass, 2 needs interactive browser confirm, 1 deterministic failure. The deterministic layer (mcpo sparse+dense + llama tool-calling) is the reliable CI signal; OWUI's interactive agent loop is confirmed in the browser (it is finicky to drive headlessly — the one non-obvious requirement is `assistant_message_id` in the **chat/completions request body**, encoded in the script).

OpenNebula MCPs — mock mode & the confirmation gate (ADR-0004)

Until live OpenNebula exists, `mcp-placement` and `mcp-host-control` run in **mock mode** (`ONE MOCK/HOST MOCK` default to 1 in compose) against canned VM↔host

data, so the agent path and the human-in-the-loop UX are testable with no cluster. Set `ONE MOCK=0/HOST MOCK=0` (and fill `.env`) at deployment.

- **Read-only tools** (`where_is_vm`, `list_vms_on_host`) are exposed to the model via `mcpo (/placement)` as tool server `server:1`.
- **The mutating op is gated.** `host-control` is bridged by `mcpo` but **deliberately NOT registered as an OWUI tool server** — the model cannot call `reboot` directly. It goes only through the OWUI native Python Tool `infra/openwebui/tools/reboot_guarded.py`, which: dry-runs to get the plan → shows it in a blocking `__event_call__` confirmation modal → executes with `confirm=True` only on human approval. The server-side allowlist (`HOST_ALLOWLIST`) + `confirm` gate still apply underneath. This is the ADR-0004 pattern: the LLM is out of the confirm loop; the human is the gate.
- `setup_openwebui.py` registers both tool servers, creates the Python Tool, and attaches all three (`server:0 rag`, `server:1 placement`, `reboot_guarded`) to the `ragfarm` preset. `__event_call__` modals require an interactive UI session (they do not fire on headless/API calls — correct for human confirmation).

The AMD → Spark migration (DONE) and what is still open

ADR-0003's central claim held up: the durable layer (Open WebUI, `mcpo`, MCP servers, `search_corpus`, Qdrant) is hardware-agnostic and **was not re-architected**. The whole engine swap touched the serving plane and almost nothing else. Keep it that way.

Done, 2026-08-03/04 (build steps 01-04):

- **Inference** — `llama.cpp/Vulkan/GGUF` → **vLLM 0.26.0 serving NVFP4 safetensors**, same OpenAI-compatible contract on `:8080`. The model moved 7B → **30B-A3B MoE**.
- **Embedder** — BGE-M3 CPU → **CUDA**, same dense+sparse contract on `:8090`. The service now refuses to start without CUDA rather than silently falling back.
- **Reranker** — Vulkan → **CUDA** `llama.cpp`, `/reranking` contract unchanged. This is the only remaining `llama.cpp` user. Query-time only — never requires re-ingest.
- **Corpus** — re-ingested from scratch on the Spark: 5 files → 183 points, dense 1024
 - sparse, alias `corpus` (ADR-0006).

Still open:

- **Networking:** the `network_mode: host` workaround is still in place and still load-bearing for the same reason as before (host services bind loopback only). It

could now be revisited, but nothing forces it — treat it as optional cleanup, and re-evaluate the 0.0.0.0 exposure + firewalling for the target network first.

- **mcpo config:** mcp-placement (where_is_vm) and mcp-host-control are already mounted in services/mcp-gateway/mcpo-config.json and run in MOCK mode. When OpenNebula is reachable (steps 05/06 unblock), set ONE MOCK=0/ HOST MOCK=0 and fill ONE_XMLRPC/ONE_AUTH per services/mcp-placement/.env.example — no mcpo-config or registration changes needed.
- **Did the bigger model fix the 7B problems?** The move to 30B was expected to resolve tool-calling discipline (the repeated-reboot miss), rambling answers, and instruction-following. **Unverified end-to-end** — tool-calling works at the API level (get_weather with parsed args), but the agent-layer behaviour is step 07. Re-test rather than assuming; see the prompt-design notes below, several of which were scaffolding for an 8B and may now be unnecessary.
- **GGUF-era scripts:** fetch-llm.sh, activate-llm.sh, llama-launch.sh still assume GGUF + --mmproj and are **unused for generation**. deploy.sh no longer calls them. A vLLM-shaped model-swap equivalent does not exist yet.
- **LLAMA_DIR is half-honoured:** deploy.sh and fetch-encoder.sh respect it, ragfarm-reranker.service and llama-launch.sh **hardcode** /home/dave/llama.cpp. Either make the units read it or drop the variable — don't leave it half-and-half.

Vision engine (Qwen3-VL family — ADR-0009)

There is one model now, and it is the vision model. Qwen3-VL-30B-A3B-Instruct is both the general and the vision engine on the Spark — there is no separate text model to flip to, and no --mmproj to attach: vision is native to the checkpoint and vLLM serves it directly (ADR-0009 as amended by ADR-0013).

To change the served model, edit LLM_MODEL_PATH / LLM_SERVED_NAME in .env and sudo systemctl restart ragfarm-vllm. Keep the alias in step with the MODEL_TUNING key in infra/openwebui/setup_openwebui.py or the preset silently stops binding.

The old activate-llm.sh --dir ... + --mmproj flow documented here was the GGUF / llama.cpp mechanism. Those scripts still exist but are **not** used for generation, and there is no vLLM-shaped replacement yet (see "Still open" above).

Live capabilities that Just Work

The Qwen3-VL preset has already been verified end-to-end on this stack:

- **OCR from image URL** — see the demo commands below. Auntie Anne's Indonesian receipt from the openlm.ai example was OCR'd correctly (all prices, all fields). Measured at 8 tok/s decode on the OLD AMD iGPU; the

Spark decodes far faster (see the performance tables below), so this is a capability record, not a current timing.

- **Image description** — attach any image in OWUI chat; the model describes content verbatim (RULE 4 of the vision prompt: no invention, verbatim OCR).
- **Diagram scan → regenerate as mermaid or draw.io** — screenshot a hand-drawn or existing diagram, ask "regenerate this as draw.io" (or "as mermaid").
- **Prompt-modify-synthesize** — attach an image, ask "same structure but add a reranker box between retrieval and generation". The model reads the input as a graph, applies your edit, and re-emits in the chosen format.

Diagram rendering

ragfarm-vision's system prompt asks the user which format they want (never routes silently, never emits both):

- **Mermaid** — the user says "mermaid". Rendered natively by OWUI as an SVG.
- **draw.io** — the user says "draw.io" / "drawio" / "editable" / "interactive" / "pan-zoom". The model emits a fenced ````html` block: a fixed wrapper plus the raw `<mxfile>` XML it authored. OWUI's HTML preview iframe loads it — pan/zoom/lightbox/layer toggle work in-chat. Everything comes from the local `drawio-viewer` nginx container, so it works air-gapped.

draw.io: the six things that must all be true

Every one of these used to fail **silently** — OWUI rendered an empty white box, with no error in the UI, the container logs, or anywhere else. Four were broken at once on 2026-08-09 and the symptom was identical in each case, so check them in this order rather than guessing. Items 5 and 6 now paint a red message into the pane instead of nothing, so in practice a truly blank pane today means 1-3.

1. **The mirror must be on disk.** `infra/drawio-viewer/` is 153 MB of the jgraph/drawio webapp, gitignored, and nothing in the build pulls it — a fresh clone has only the two tracked HTML files, and nginx 404s `js/viewer-static.min.js`. Fix: `scripts/fetch-drawio-viewer.sh` (idempotent, no-ops when present). Step 07's deploy fragment runs it.
2. **The URLs must name a host the CLIENT can reach.** This is the one place in the whole system where `127.0.0.1` is wrong: the wrapper is fetched by the user's browser, so on a remote client it means that laptop's loopback. `setup_openwebui.py` bakes `RAGFARM_PUBLIC_HOST` into the wrapper, falling back to autodetecting the default-route address.

3. **IFRAME_CSP must whitelist that same host.** OWUI's preview iframe otherwise refuses the script. `infra/compose.yaml` reads `${RAGFARM_PUBLIC_HOST:-127.0.0.1}` — and compose, unlike the Python, cannot autodetect. Leave the var unset and you get the worst combination: a correctly autodetected URL that the CSP then blocks. So **set it in `.env`** `setup_openwebui.py` prints a warning when it is missing. Changing it needs both `docker compose up -d open-webui` (CSP is read at container start) and a `setup_openwebui.py` re-run.
4. **The XML must reach the viewer the way the viewer expects.** It does not read a child `<xml>` element — that form throws can't access property "length", a is undefined and draws nothing. The XML goes in the JSON `data-mxgraph` attribute under an `xml` key. All of that now lives in `infra/drawio-viewer/ragfarm-drawio.js`, which the wrapper's last line loads; the model writes raw XML into a `<script type="application/xml">` tag and never has to escape anything.

That file is tracked in git even though the rest of `infra/drawio-viewer/` is not, and it exists for a specific reason: **boilerplate the model must retype is a liability proportional to its length.** The wiring used to be 15 lines of inline bootstrap in the wrapper, and on a 31-node diagram the model reproduced the whole page correctly except those lines — emitting a valid `<mxfile>` that pasted into draw.io online perfectly, and a blank pane in chat. It did keep the one-line `<script src>` right after them. So the logic moved into the file that one line loads. Fixture: `tests/fixtures/splunk-kb-model-output.drawio` is that answer's XML.
5. **The reply must be inside a ````html` fence.** OWUI only turns a fenced html block into a diagram pane; raw HTML in the message body renders as nothing. The model gets this right on small diagrams and has been observed dropping it on a 24-entity one, after 9k tokens of reasoning. RULE 5 now demands the fence as the first characters of the reply.
6. **The reply must not be cut off by `max_tokens`.** Reasoning and answer share one budget (see the sizing note in `MODEL_TUNING["vision-thinking"]`), and a diagram costs 9.5-13.3k completion tokens of which roughly two thirds is reasoning. `ragfarm-drawio.js` prints "the answer did not finish" into the pane when `</mxfile>` is absent, so this no longer presents as a blank box.

Smoke test, before blaming the model: open `http://<RAGFARM_PUBLIC_HOST>/fixtures/drawio-wrapper-reference.html` in the demo browser. It is the wrapper verbatim, with a known-good two-box diagram. Renders → items 1-3 are fine

and the blank pane came from the model's XML. Blank → it is the infrastructure, not the model.

Model-side XML failures (RULE 5 spells these out, all seen in practice): eliding attributes as ... /> instead of writing them out; reusing the reserved id="0" / id="1" for content cells; omitting <mxGeometry>, so a cell has no position or size; and re-deriving a user-supplied diagram from scratch instead of editing their XML in place, which throws away their layout and colours.

Verified demo commands (no prep needed, run today)

The tests below use the live stack as-is. Nothing to fetch, nothing to install.

Sanity: which model is served?

```
curl -s 127.0.0.1:8080/v1/models | python3 -c 'import sys,json; print([m["id"]
for m in json.load(sys.stdin)["data"]])'
# expect: ['qwen3-vl-30b-a3b']
```

(vLLM returns the OpenAI {"data":[...]} shape. llama.cpp's old {"models":[...]} with a capabilities list is gone — a command reading d["models"][0] is pre-Spark.)

OCR from a public image. Sent as base64 rather than a URL — originally to work around llama-server's HTTPS-off build, and still the safer habit on an offline box:

```
.venv/bin/python - <<'PY'
import base64, requests, time
IMG = "https://ofasys-multimodal-wlcb-3-toshanghai.oss-cn-shanghai.aliyuncs.com/
wpf272043/keepme/image/receipt.png"
img = requests.get(IMG, timeout=30); img.raise_for_status()
uri = f"data:{img.headers['content-type']};base64,
{base64.b64encode(img.content).decode()}"
t0 = time.time()
r = requests.post("http://127.0.0.1:8080/v1/chat/completions", json={
    "model": "qwen3-vl-30b-a3b",
    "messages": [{"role":"user","content":[
        {"type":"image_url","image_url":{"url":uri}},
        {"type":"text","text":"Read all the text in the image."}]}],
    "max_tokens": 512, "temperature": 0.6,
}, timeout=120)
print(f"HTTP {r.status_code} {time.time()-t0:.1f}s")
print(r.json()["choices"][0]["message"]["content"])
PY
```

In-UI vision demo (OWUI, admin@ragfarm.local):

1. Select model ragfarm-vision from the dropdown.

2. Click the paperclip → upload any image (receipt, screenshot of a diagram, photo of a whiteboard).
3. Ask one of: "OCR everything in this image", "describe what you see", "regenerate this diagram as mermaid" / "as draw.io".
4. For a diagram request, add `/no_think` to the prompt if the Thinking trace makes the wait too long (Qwen3 chat-template convention: strips the `<think>...</think>` block for that turn).

Serve local test images to the model: the same `drawio-viewer` `nginx` also serves anything under `infra/drawio-viewer/`. Drop a `.png` / `.pdf` there and reference `http://127.0.0.1:8091/<file>` in a prompt. Useful for a repeatable demo without leaning on external URLs.

`/think` vs `/no_think` (Qwen3 Thinking control)

Applicability on the Spark: the deployed checkpoint is `Qwen3-VL-30B-A3B-**Instruct**`, not a `*-Thinking-*` variant, so there is no reasoning block to suppress and `/no_think` is a no-op for it. Kept because the convention is real, survives a model swap to a Thinking variant, and the demo advice below still applies if one is ever served.

Qwen3 parses these tokens out of user messages (chat-template convention, NOT system-prompt rules):

- `/think` — force the model to emit a `<think>...</think>` block before the answer. Default for `*-Thinking-*` model variants.
- `/no_think` — suppress the reasoning block for this turn. Runs a Thinking model like an Instruct one; 3-5× snappier answers, at the cost of the reasoning quality on hard multi-step prompts.

Practical demo advice: default (thinking on) for the first, hardest question of the day (RAG lookup, complex OCR); append `/no_think` for follow-ups, quick lookups, and diagram requests where the trace adds no value.

Debug & measurement (tests/tracing/)

Standalone Python tools (no dependencies beyond `requests`) that answer where did the time go for any inference or chat turn. All safe to run against the live stack — read-only, no side effects.

Status: incomplete, rewrite expected. This is a framework, not a finished tool, and it has no concept of thinking/reasoning models. Some docstrings still quote pre-Spark ports. Treat output as indicative; requirements for the rewrite are in `tests/tracing/README.md`. The `owui_trace_proxy.py` experiment (port 8095) is **retired** — it

corrupted the OWUI DB and presets during demo prep. It is still on disk; do not re-plumb it without a plan.

Every tool takes `--url http://127.0.0.1:8080` (or the equivalent flag) so they're portable across the swappable llama-server endpoint. Default endpoint in the code is `localhost:8001` — always pass `--url` explicitly.

The catalog

Tool	What it measures	When to use
<code>ragfarm_bench.py</code>	Basic prefill/decode tok/s + TTFT + E2E per prompt	Quick "is llama fast enough right now" check
<code>ragfarm_bench_extended.py</code>	Full per-stage timings, absolute token+byte counts, context growth, CSV/JSON export	Regression tracking across commits or hardware swaps
<code>ragfarm_bench_chatid.py</code>	Same, plus context-blowup detection per chat session	Find when a conversation starts overflowing context
<code>chat_execution_tracer.py</code>	Chat session timeline: user→prefill→tool decision→tool exec→decision phase→generation, with orchestration-overhead %	Diagnose "chat feels slow" vs "LLM is slow"
<code>ragfarm_tracer_simple.py</code>	One-shot telemetry query: which model is loaded, endpoint latencies	Sanity check before any deeper trace
<code>ragfarm_integrated_tracer.py</code>	Combined engine telemetry + pipeline trace demo	Report generation for a whole pipeline
<code>ragfarm_rag_tracer.py</code>	RAG candidate-pool evolution: Qdrant → RRF → reranker, per-stage tokens/latency	Debug retrieval quality regressions ("why didn't my chunk win?")

(Not integrated: *ragfarm_http_tracer.py* — proxy-based, requires reconfiguring OWUI's LLM endpoint. Its only unique benefit is E2E prompt-answer wall time, which every other tool measures anyway.)

Verified one-liners (run today)

```
# 1. What model is loaded and how fast does one prompt run?
.venv/bin/python tests/tracing/ragfarm_tracer_simple.py query \
  --generation localhost:8080 --reranker localhost:8081

# 2. Baseline bench, one prompt, small max_tokens (fast)
.venv/bin/python tests/tracing/ragfarm_bench.py \
  --url http://127.0.0.1:8080 --prompt 1 --max-tokens 40

# 3. Extended bench, 3 prompts with CSV export (regression file)
.venv/bin/python tests/tracing/ragfarm_bench_extended.py \
  --url http://127.0.0.1:8080 --prompt 3 --max-tokens 128 \
  --csv bench_$(date +%s).csv

# 4. Chat tracer demo (canned session; no LLM required)
.venv/bin/python tests/tracing/chat_execution_tracer.py --demo
# writes chat_trace_demo.json in cwd

# 5. Real RAG pipeline trace for one query (endpoint is mcpo at :8000,
#   which mounts rag under /rag/search_corpus — NOT direct rag-retrieval :8104)
.venv/bin/python tests/tracing/ragfarm_rag_tracer.py trace \
  --chat-id demo_$(date +%s) \
  --query "FW pravidla pro host leadb229p.lea.piz" \
  --rag-endpoint http://127.0.0.1:8000
```

What "good" looks like — Spark (current) vs AMD iGPU (historical)
Current: DGX Spark GB10, Qwen3-VL-30B-A3B NVFP4 on vLLM 0.26.0.

Metric	Now	Comment
Decode	75.6 tok/s (flashinfer_b12x) / 71.1 tok/s (FLASHINFER_CUTLASS, default)	256 tok, ignore_eos, single stream, 3 runs. 54% of the 140 tok/s bandwidth ceiling
Prefill	UNMEASURED	Where FP4 should show its biggest win — the gap most worth closing
Vision OCR (dense receipt)	UNMEASURED on Spark	Capability verified on AMD; re-time it here

Reranker turn	UNMEASURED on Spark	Was 1.9 s on the iGPU; same model, now CUDA
Tool overhead	UNMEASURED on Spark	Was 15–25% of a chat turn
Cold start (vLLM)	813 s	FlashInfer JIT; torch.compile is only 9.5 s of it
Warm start (vLLM)	3 min	Reuses <code>~/.cache/flashinfer</code> + <code>~/.cache/vllm</code>

Most rows are still UNMEASURED because step 07 (agent wiring) is where chat-turn and tool timings become meaningful. Fill them in from real runs — do not carry the AMD numbers across.

Historical baseline: AMD Ryzen MiniPC iGPU, Qwen3-VL-8B-Thinking Q4_K_M / Vulkan. Kept deliberately as the before/after reference — this is what the hardware investment was measured against.

Metric	AMD iGPU	Comment
Decode	8–9 tok/s	LPDDR5x bandwidth-bound; both 7B and 8B landed here
Prefill	170 tok/s	The order-of-magnitude gap vs decode is expected
Vision OCR (dense receipt)	40 s	300+ output tokens at 8 tok/s
Reranker turn	1.9 s	After ADR-0008 moved it to GPU/Vulkan (was 36 s on CPU)
Tool overhead	15–25 % of chat turn	Higher with the <code>reboot_guarded</code> modal, lower for pure RAG

Headline so far: decode went 8.5 → 71 tok/s, roughly 8×, while the model grew from 8B dense to 30B-A3B MoE — i.e. more capable and an order of magnitude faster. Note the comparison is not like-for-like (different

model, quant and engine); it is the honest end-to-end "what the box does for us" delta, which is the number that justified the hardware.

Older docs citing 300 tok/s prefill / 1000 tok/s decode refer to a much lighter model than either row above — ignore them.

Slots — two models resident, switchable mid-chat

Measured 2026-08-05 with both slots live:

slot	unit	port	model	util	GPU
0	ragfarm-vllm@0	8080	Qwen3-VL-30B-A3B-Thinking NVFP4 (MoE)	0.270	30.1 GB
1	ragfarm-vllm@1	8082	Qwen3-VL-32B-Thinking NVFP4 (dense)	0.291	35.4 GB

Plus embedder 1.6 GB and reranker 0.9 GB: **68 GB of 121 GB**, 33 GB free. The derived budget predicted 0.561 and the box landed on it.

The formula is (weights + 12 GiB KV + 3 GiB overhead) / 121.7, and vLLM's own startup report confirms it is sound — slot 0 was granted 32.86 GiB and used 18.22 weights + 5 overhead + **9.91 KV**, i.e. the 12 GiB KV allowance is slightly generous, which is the right direction to be wrong in.

Open WebUI needs the endpoints registered through its API, not just compose. OWUI seeds `OPENAI_API_BASE_URLS` from the environment on FIRST start only and then persists it in its own DB. Changing compose on an existing deployment leaves it on one endpoint and the second model silently never appears in the model list. `setup_openwebui.py` now POSTs `/openai/config/update` with every slot URL, so a re-run fixes it; that is also why the script must be re-run after adding a slot.

Cold start on a new checkpoint (JIT cache partially missing): slot 0 took 3.5 min, slot 1 7 min. Warm restarts are far quicker.

Chapter 6 LLM life-cycle — models, slots, running the stack

LLM life-cycle — models, slots, and running the stack

How a model gets from Hugging Face onto the box, into GPU memory, and in front of a user; and how the whole system is started, stopped and checked. Three tools and one registry, each with a full manual page under `docs/man1/`.

```
fetch_llm.py      -> on disk + registered      "we have it"
activate_llm.py   -> bound to a slot, serving "it is resident"
setup_openwebui   -> preset bound to its alias "a human can pick it"
stack.sh          -> the whole system up, and honestly checked
```

What a slot is

vLLM serves exactly **one** base model per process. There is no multi-model mode: several `--served-model-name` values are aliases for the same weights, and LoRA adapters need a shared base. A slot is therefore this project's name for one vLLM instance:

slot	systemd unit	endpoint
0	ragfarm-vllm@0.service	127.0.0.1:8080/v1
1	ragfarm-vllm@1.service	127.0.0.1:8082/v1

Port is $8080 + 2N$; the stride is 2 because 8081 belongs to the reranker.

Why two. Two occupied slots mean two models resident at once, and Open WebUI can switch between them **mid-conversation with the context intact**. That is not a convenience feature — it is the differential-diagnosis instrument. When an answer is wrong, asking a second model the same question with the same history separates "the model cannot do this" from "our prompt told it not to". Most of the August prompt fixes were found that way.

The memory budget, which is the whole difficulty

`--gpu-memory-utilization` is a fraction of **total** device memory, and every vLLM instance computes it **independently** — neither slot knows the other exists. Two slots each asking for the single-model default of 0.50 will OOM. On a GB10 there is no separate VRAM to retreat into: 121.7 GiB is shared with the embedder, the reranker, Qdrant, the container plane and the operating system.

So the per-slot fraction is derived, never chosen:

```
util = (weights + KV_TARGET_GIB + OVERHEAD_GIB) / TOTAL_GIB
```

with `weights` taken from the size recorded **after** the download was verified against the Hub — never from `du`, because a half-finished checkpoint measures small and would under-allocate the slot. The sum across slots must stay under `BUDGET_CEILING` (0.72), leaving at least 28% for everything that is not an LLM. Exceeding it is a hard refusal, not a warning.

Today's configuration, as reported by `activate_llm.py --status`:

slot	port	util	model
0	8080	0.338	Qwen3-VL-30B-A3B-Thinking-FP8 [MoE]
1	8082	0.237	Qwen3-VL-30B-A3B-Instruct-NVFP4 [MoE]

total GPU budget: 0.575 of 0.72 ceiling

The registry

`models/llm/active.json` is git-tracked and answers two questions that change on different timescales:

- `downloaded[]` — every model this deployment should have on disk. You fetch a model once.
- `active[]` — which of them are resident, one array position per slot, as **indexes into** `downloaded[]`. You re-bind slots several times a day.

Storing the binding as an index means a slot cannot reference a model that was never registered. Each entry carries the directory name, the Hub repo, the served alias, the Open WebUI preset id, a display label, a tuning profile and the verified size. A machine can be rebuilt from this file with one command.

`model`, `alias` and `preset` must each be unique. The tools **refuse to run** on a duplicate rather than disambiguating, because choosing which name wins is a decision only a person should make — silently renaming produces a model picker whose entries do not mean what they say.

Everyday commands

```
scripts/stack.sh status           # every service, endpoint, state
scripts/activate_llm.py --status  # slots, models, memory budget
scripts/activate_llm.py           # interactive: activate or clear
scripts/activate_llm.py -s 0 -m Qwen3-VL-30B-A3B-Thinking-FP8
```

One call rewrites the slot's config, restarts its unit, waits for it, and re-binds the Open WebUI presets.

Adding a model:

```
scripts/fetch_llm.py -m <hf-repo> \
  --alias <served-name> --preset <owui-preset-id> \
  --display '<label for the picker>' --profile vision-instruct
```


Uniqueness is checked **before** the transfer, so a name clash costs a typo rather than an afternoon and 70 GB.

Four things that will bite

A model swap without re-binding the presets. Presets bind by alias. Change what a slot serves and the preset names an alias nothing serves; the picker then offers only raw base models — no system prompt, no tools, no grounding rules. Nothing looks broken and the answers are silently ungrounded. `activate_llm.py` now re-binds automatically, and prunes presets whose model is gone.

Starting slots in parallel. vLLM profiles free GPU memory during initialisation; if a second instance is allocating at the same moment the first sees memory move underneath it and dies with `Error in memory profiling OR No available memory for the cache blocks`. Neither message names the race. The tools serialise; `systemctl start a b c` does not.

The reasoning parser is a property of the checkpoint. A Thinking model needs `--reasoning-parser qwen3`; an Instruct model must not have it. The qwen3 parser keys on `</think>`, and finding none it files the model's entire answer as reasoning — the chat message comes back empty with the whole reply inside the collapsed thinking trace, which is convincingly mistakable for a broken quantisation. It is written per slot from the registry's `profile` field.

max_tokens shares the context window with the prompt. vLLM does not clamp it; it rejects the request. A generous cap works until a tool result grows the prompt on the continuation request, and then every tool-using turn returns nothing. With `--max-model-len 32768`, `max_tokens 8192` leaves 24 k of prompt headroom, which is the reason it is set there.

Boot behaviour

Slot 0 is enabled at boot; slot 1 deliberately is not. Enabling both would have systemd start them in parallel, which is exactly the race above, and the serialisation lives in the tooling rather than in the unit. The box comes up serving the MoE; a second model is a deliberate act.

After a reboot, this is the check:

```
scripts/stack.sh status
```

It probes all thirteen services and reports `[OK]`, `[NOT_OK: code]`, `[DEGRADED]`, `[OFF]` or `[ABSENT]`. `[DEGRADED]` is the interesting one — a service that answers correctly and is still broken, such as `mcpo` serving 200 with zero tools mounted, or `nginx`

serving its landing page while the draw.io viewer JavaScript is absent. Both of those happened, and both were invisible to a status-code check.

Manual pages

Full detail, including the single-large-model workflow and the failure modes above with their diagnostics:

```
man docs/man1/stack.1          # lifecycle and health
man docs/man1/activate_llm.1   # slots and the memory budget
man docs/man1/fetch_llm.1      # downloading and registering
man docs/man1/active.json.1    # the registry format
man docs/man1/setup_openwebui.1
man docs/man1/env.1            # how configuration propagates
```

Chapter 7 Prompt library and regression suite

Prompt library — demo reference and regression suite

This file has two jobs and one format.

For a human it is the demo script: what to ask, what a good answer looks like, and which of those answers has actually been seen rather than hoped for.

For a machine it is the regression suite. `tests/promptlib.py` parses every case out of it, `scripts/test_regressions.py` replays them against the live slot and an LLM judge decides whether today's answer still achieves what the `expect` block describes. One file, so the demo script and the tests cannot drift apart.

How to add a case

Copy the shape below. Only the heading, `~~~prompt` and `~~~expect` are required.

```
### X1 · short title
```

```
- **tags:** rag, czech  
- **tools:** search_corpus  
- **status:** unverified
```

```
~~~prompt
```

```
The prompt, verbatim, exactly as a user would type it.
```

```
~~~
```

```
~~~expect
```

```
What a correct answer must achieve, in prose.
```

```
~~~
```

```
~~~must
```

```
a-substring-that-must-appear
```

```
~~~
```

Rules worth knowing before you write one:

- **Fences are tildes**, not backticks, so a prompt or expectation can itself contain ````` blocks. The `draw.io` and `code` cases do.
- **expect is prose, never a golden string**. Sampling is non-greedy, so two correct answers differ in wording every run. Describe what the answer must achieve the judge decides whether it did. A literal diff would fail every time and everyone would learn to ignore the suite.
- **must / must-not are for the deterministic part** — a hostname, an identifier, a figure. They are checked by plain string match before any model is called, because a cheap check that can fail hard should not be delegated to a judge.

- **status: is a claim about evidence.** Say which model and which date, or say unverified. Several cases below are verified only on models we no longer run; that is recorded rather than hidden.
- **skip:** with a reason keeps a case in the library but out of the run.

Validate after editing:

```
.venv/bin/python tests/promptlib.py --validate
```

Metadata keys: preset, slot, tools, attach, status, tags, skip.

Section R — Retrieval and grounding

The core of the product. Every case here must call `search_corpus` and answer only from what comes back.

Provenance warning. R1-R6 were verified on 2026-07-27 against the retired text preset (Qwen2.5-7B Q4_K_M, greedy) and have README screenshots. They were not re-verified on the current 30B models except where the status line says so. That is exactly the drift this suite exists to catch.

R1 · FW rules for one specific host

- **tags:** rag, czech, table
- **tools:** search_corpus
- **status:** verified 2026-07-27 (text preset, screenshot ex-01); tool call re-verified 4/4 on 30B-A3B Instruct 2026-08-12

Jaká jsou FW pravidla pro host leadb229p.lea.piz?

Calls `search_corpus`, then returns the firewall rules for that one host as a markdown table with every column present: source and destination network addresses, source and destination network names, destination ports, protocol. Around six rows. Ends with a Source: line naming the xlsx it came from. Does not dump rules for other hosts.

leadb229p

R2 · Contacts for project management

- **tags:** rag, czech, list
- **tools:** search_corpus
- **status:** verified 2026-07-27 (text preset, screenshot ex-02)

Dej mi kontakty na projektové vedení v EPC.

Calls `search_corpus` and returns six people as a list, each with company, role/area, phone and e-mail. A phone number is omitted only where it is genuinely absent from the source. No invented contacts.

R3 · Everything about one host

- **tags:** rag, czech, table
- **tools:** search_corpus
- **status:** verified 2026-07-27 (text preset, screenshot ex-03)

Co vis o hostu acclclass1?

Calls search_corpus and returns the full record for that host – on the order of nineteen fields, including environment, OS, virtualisation, storage, vCPU, RAM, HDD, support, hostnames, IP addresses, netmasks, VLAN id, filesystem and UID. For this one host only.

acclclass1

R4 · Where credentials live

- **tags:** rag, czech
- **tools:** search_corpus
- **status:** verified 2026-07-27 (text preset, screenshot ex-04)

Kde ukládáme hesla pro EPC?

Calls search_corpus and explains that passwords are kept in NordPass, names the record to look up, and says how to use it with the RDP client. Does not print any actual credential.

password:

R5 · Where the documentation repo is

- **tags:** rag, czech
- **tools:** search_corpus
- **status:** verified 2026-07-27 (text preset, screenshot ex-05)

Kde máme uložené GIT repo s dokumentací?

Calls search_corpus and returns the exact Azure DevOps URL for the sa-hosting repository, not a paraphrase of it.

dev.azure.com

R6 · How to log in from ŠA to EPC

- **tags:** rag, czech, procedure
- **tools:** search_corpus
- **status:** verified 2026-07-27 (text preset, screenshot ex-06)

Jak se přihlásím ze ŠA do EPC?

Calls search_corpus and lays out the login paths as a numbered procedure – SSH direct, RDP direct, terminal servers, GUI, reverse proxy, CLI – with hostnames, IP addresses, where the credentials come from, and which spreadsheets hold the details.

R7 · All CutOver contacts

- **tags:** rag, czech, list
- **tools:** search_corpus
- **status:** tool call verified 4/4 on 30B-A3B Instruct and 30B-A3B Thinking, 2026-08-12

Úplně všechny kontakty co máme v souvislosti s CutOver plánem?

Calls search_corpus and returns every person listed on the CutOver contact sheet, with role and contact details. This is an overview question, so the full table is the right answer rather than a single record.

CutOver

R8 · General knowledge is answered directly

- **tags:** general, czech, negative
- **status:** rewritten 2026-08-14 — the July expectation (route everything through search_corpus) contradicts today's RULE 0 carve-out

Kdo je prezident USA?

Answers directly from the model's own knowledge and does NOT call search_corpus: RULE 0 makes general-knowledge questions tool-forbidden, and the corpus has nothing to say about US presidents. Names the president its training data knows (Joe Biden for the Qwen3-VL generation) and may add that its knowledge has a cutoff. What must not happen is a corpus lookup, a claim that the corpus contains the answer, or a refusal to answer at all.

R9 · Honest miss inside the corpus domain

- **tags:** rag, czech, negative
- **tools:** search_corpus
- **status:** new 2026-08-14, unverified — replaces the anti-hallucination half of the old R8

Co víš o hostu zzql-nonexistent-42?

This IS a corpus question, so it calls search_corpus. Finding nothing for that host, it says plainly that the corpus contains no such host. It must not invent an IP address, VLAN, owner or any other field, and must not silently answer about a different host whose name merely looks similar.

Section T — Tools beyond retrieval

T1 · Clock

- **tags:** tools
- **tools:** get_current_timestamp
- **status:** verified 2026-07-27 (text preset, screenshot ex-07)

Kolik je hodin?

Calls the clock tool and answers with one short line giving the current time and the timezone offset. Does not call search_corpus.

T2 · Human-gated reboot

- **tags:** tools, safety
- **tools:** reboot_host
- **status:** verified 2026-07-27 (text preset, MOCK placement backend, screen-shot ex-07)
- **skip:** mutating action behind a confirmation modal; cannot run unattended

Rebootuj host node-03.

Routes to reboot_host through the reboot_guarded confirmation wrapper and STOPS for human approval. On approval, reports that the host was restarted and how many VMs were drained first. It must never report success without the confirmation step having happened.

Section V — Vision

Provenance warning. V1-V7 were verified on 2026-07-27 against Qwen3-VL-**8B**-Thinking (GGUF, old AMD box). Timings quoted in the status lines are from that hardware and are historical; the Spark is far faster. The capability claims are what matter.

V1 · Receipt OCR, real world

- **tags:** vision, ocr
- **status:** verified 2026-07-27 on 8B-Thinking — 77 s, 613 out tokens, 8.1 tok/s (old AMD iGPU)

Read all the text in this image. Preserve the original numbers and layout.

Transcribes every line of the receipt verbatim, including the vendor name, the line item with its quantity and unit price, subtotal, grand total, cash tendered and change due. All figures exact. Nothing invented, nothing translated unless asked.

V2 · Mixed-script OCR with identifiers

- **tags:** vision, ocr
- **status:** verified 2026-07-27 on 8B-Thinking — 59 s, 462 out tokens

OCR every visible line. Keep punctuation, numbers, and IP addresses exact.

Every IP address, CIDR suffix, port, e-mail address and float reproduced character for character, including negative and multi-decimal values. Czech words kept as written. Zero hallucinated lines.

V3 · Chinese to English

- **tags:** vision, ocr, translation
- **status:** verified 2026-07-27 on 8B-Thinking — 50 s, 394 out tokens

This image contains Chinese text. Read each line and give an English translation.

Reads each line of Simplified Chinese and gives a correct English translation, preserving the numbers and units exactly – server counts, bandwidth figures and the maintenance window with its times.

V4 · Farsi to English

- **tags:** vision, ocr, translation
- **status:** verified 2026-07-27 on 8B-Thinking — 190 s, 1500 out tokens; slowest case in the library

This image contains Persian (Farsi) text. Read every line and provide an English translation next to each.

Recognises the Persian text, breaks it into words and translates each line into English alongside the original. Place names and temperature units come through correctly. Digit misreads on non-Arabic-shaping fonts are a known limit and not a failure of this case.

V5 · Bar chart to table

- **tags:** vision, extraction, table
- **status:** verified 2026-07-27 on 8B-Thinking — 28 s, 211 out tokens; every value correct

This image is a bar chart. Extract the underlying data into a Markdown table with two columns: Region and Revenue.

A markdown table with exactly the two requested columns and one row per bar, the region labels transcribed and each revenue value read correctly from the bar height. No commentary needed beyond the table.

V6 · Diagram photo to mermaid

- **tags:** vision, diagram
- **status:** verified 2026-07-27 on 8B-Thinking — 134 s, 1061 out tokens

This image shows a small architecture diagram. Regenerate it as a mermaid graph (fenced mermaid). Preserve every box and every arrow direction.

One fenced mermaid block. Every box from the source appears as a node, every arrow appears with its direction preserved, and no edges are invented. Open WebUI renders it inline.

mermaid

V7 · Code photo to runnable Python

- **tags:** vision, code

- **status:** verified 2026-07-27 on 8B-Thinking — 155 s, 1232 out tokens; transcribed $n^{**0.5}$ as $n*0.5$, still predicted the right output

This image shows Python code. Transcribe it verbatim into a fenced python block. Then predict what it prints (do not use any tool).

A fenced python block reproducing the code from the photo, followed by a prediction of its output. The predicted output must be correct. It must NOT call the code interpreter, because the prompt forbids it.

Section D — Diagrams the model authors

D1 · Reverse the arrows in a supplied diagram

- **tags:** diagram, drawio, edit
- **attach:** tests/fixtures/dependency_input.drawio
- **status:** verified 2026-08-09, 10/10 structural checks on both 30B-A3B Thinking FP8 and 32B dense FP8; rendered headlessly

V přiloženém Draw.io grafu obrať směr šípek mezi Frontend → API → Database opačným směrem. Prezentuj výstup opět ve formátu Draw.io.

One fenced html block containing the ragfarm draw.io wrapper and the user's own XML with only the arrow directions changed: each edge's source and target swapped. The user's cell ids, styles, geometry, labels and fill colours are preserved exactly – the diagram must be edited in place, not redrawn. No elided attributes, no reserved ids 0 or 1 on content cells, every cell keeps its mxGeometry, and the page loads ragfarm-drawio.js.

```
ragfarm-drawio.js
```

...

D2 · Author a diagram from a description

- **tags:** diagram, drawio
- **status:** verified 2026-08-09 on 30B-A3B Thinking FP8; fenced 3/3, 11k completion tokens

Vytvoř Draw.io ER diagram naší monitorovací domény: entity Incident, Problem, OBS Process, Micro Service, Cluster, Node, Data Centre, Squad, Person. Každá entita má 2-3 atributy, barevné výplně podle domény a ortogonální spoje. Prezentuj jako Draw.io.

One fenced html block using the wrapper. Every named entity present, each with its attribute rows, colours grouped by domain, and orthogonal edges between related entities. Valid draw.io XML that renders.

```
ragfarm-drawio.js
```

D3 · 1:1 conversion of a diagram image

- **tags:** diagram, drawio, vision, hard

- **attach:** tests/fixtures/Splunk_in_KB.png
- **status:** KNOWN HARD — single-pass fails. See docs/measurements/2026-08-09-instruct-vs-thinking-drawio.json
- **skip:** single-pass conversion is a known failure; kept as the reference for the two-pass workflow

Převeď obrázek grafu z přílohy do formátu Draw.io. Zachovej formátování, rámečky, spoje, šipky, typy a styly šipek, velikosti, poměry, vztahy, barvy, zkrátka vše 1:1. Prezntuj výsledný Draw.io graf.

Twenty-eight entities as swimlanes with their attribute rows, the PoC container, the Legenda frame, colours by domain, and roughly twenty-six edges. In one pass neither model achieves this: Thinking stops early with a simplified 43-box diagram, Instruct loops on edges until the budget is gone. The working route is two passes – boxes first with edges forbidden, then edges given the box ids.

D4 · Inventory a diagram image without drawing it

- **tags:** vision, extraction, table
- **attach:** tests/fixtures/Splunk_in_KB.png
- **status:** verified 2026-08-09 on 30B-A3B Thinking FP8 — 28/28 entities, all attribute rows correct, one placement error

Vypiš úplný inventář tohoto ER diagramu jako markdown tabulku. Jeden řádek na entitu: název | barva výplně | seznam VŠECH atributových řádků uvnitř rámečku | je uvnitř modrého rámu PoC? (ano/ne). Nic nevynechávej, nic neshrnuj.

A markdown table with one row per entity – twenty-eight of them – listing every attribute row inside each box, its fill colour, and whether it sits inside the PoC frame. Source typos may be normalised. This is the perception half of D3 and it is reliable where the drawing half is not.

Section C — Coding

C1 · Write and actually run Python

- **tags:** code, interpreter
- **status:** verified 2026-07-27 (text preset, screenshot ex-09)

Vygeneruj mi kód pro quicksort a otestuj ho spuštěním nad malým polem náhodných řetězců. Prezntuj kód a výsledné pořadí tříděného pole po běhu sortu.

Emits complete Python implementing quicksort over a random list of strings, then CALLS the code interpreter to run it and reports the sorted result. Per RULE 6, Python is the one language it must execute rather than only describe.

C2 · Non-Python code is not executed

- **tags:** code, negative
- **status:** verified 2026-08-05 as the fix for the Python/JavaScript reasoning loop

Napiš mi v JavaScriptu funkci, která z pole objektů udělá mapu podle klíče id.

Outputs the JavaScript and stops. It must NOT call the code interpreter, must not invent test cases for it, and must not apologise for being unable to run JavaScript – the interpreter is Python-only and declining to run other languages is correct behaviour, not a limitation worth narrating.

code interpreter

Section M · Mermaid

M1 · Dependency tree of a sentence

- **tags:** diagram, mermaid
- **status:** verified 2026-07-27 (text preset, screenshot ex-08)

Vygeneruj stromový diagram slovních vazeb ve větě: "Once upon a time there was a very little dog called Steven who owned a nice little yellow car".

One fenced mermaid block containing a genuine branching dependency tree rather than a linear chain: the head noun acts as a hub with its determiners and modifiers as children, and the relative clause branches off correctly.

mermaid

Chapter 8 Prompt regression testing

Prompt regression testing

Why this exists

The system prompt is the entire behavioural specification of this assistant. Everything the model does that we care about — when it reaches for a tool, how much it retrieves, whether it answers with a table or a sentence, whether it admits a miss instead of inventing — is defined in prose in one string.

That string is **fidgety**. A change that reads as local to RULE 3 shifts how RULE 1 is obeyed. Every prompt defect found in August was of that shape:

- "call the tool at most once, do not refine" was read as be minimal, and the model set $k=1$, starved its own retrieval and answered from one wrong row
- "show records as a full table" outranked the user's "only him", so a question about one person returned the whole team
- a coding rule demanding interpreter verification made the model loop on "the user wants JavaScript, which is not Python. Wait, this is a problem."

None of those were model defects. All of them were found by a human noticing a bad answer, usually late. This suite makes the blast radius of a prompt edit visible in one command, before a demo does it for you.

The library is the suite

`docs/prompts.md` is the single source. It is the human demo script and the test corpus, because two copies would drift. Each case carries the prompt verbatim and an `expect` block describing what a correct answer must achieve. See the top of that file for how to add one.

```
.venv/bin/python tests/promptlib.py --validate    # lint the library
.venv/bin/python tests/promptlib.py --list       # ids, titles, tags
```

Running it

```
source scripts/proxy-env.sh
.venv/bin/python scripts/test_regressions.py      # everything
.venv/bin/python scripts/test_regressions.py --tag rag # one slice
.venv/bin/python scripts/test_regressions.py --id R1 --id R7 # named cases
```

Exit status is the worst verdict seen, so it composes into other scripts:

RC	verdict	meaning
0	EQUAL	every case achieves what it should
1	DRIFTING	something slipped; still useful answers

2	WRONG	at least one case failed outright
3	—	harness error: stack down, library unparsable

Every run writes `logs/regressions-<UTC>.json` (the record to diff between prompt versions — that is the whole point) and a matching `.log` with the full prompts, answers and timings.

Two phases, and why they are separate

Collect replays each prompt against the live slot, mimicking Open WebUI as closely as possible: the preset's own system prompt, the preset's own sampler, the same mcpo tools. The expectations in `prompts.md` were written from OWUI chats, so the replay has to come from the same shape of environment or the comparison is unfair. The preset and sampler are read out of `setup_openwebui.py` rather than restated, because a second copy of the sampler values is exactly how a harness drifts from the deployment it is meant to test.

Judge compares each answer to the case's `expect` with a second model call, a comparator system prompt, and **no tools at all**. That last part is deliberate. The judge's job is to reason over two given texts. Give it retrieval and it goes and checks for itself, then grades the answer against what it found rather than against `expect` — which quietly turns a regression suite into a second opinion.

Deterministic checks run **before** the judge: `must` / `must-not` substrings, empty answers, and whether the expected tool was called. A check that can fail hard for free should never be delegated to a model.

Verdicts

- **EQUAL** — achieves everything `expect` describes. Wording, ordering and formatting may differ completely; sampling is non-greedy, so they always do.
- **DRIFTING** — substantially right, but a required column, field or list item is missing, or the format slipped, or unrequested commentary buried the answer.
- **WRONG** — contradicts `expect`, is empty or truncated, refuses something it should do, answers from general knowledge where a lookup was required, invents a record, or narrates a tool call instead of making one.

Calibrating the judge

An LLM judge is a measuring instrument and needs checking like one. Two habits:

Read the reasons, not just the verdicts. The judge writes at most three sentences saying what is missing. On the first real run it returned:

R1 DRIFTING — ACTUAL includes extra columns not specified in EXPECTED. The table has only five rows instead of "around six".

The row count is a fair catch. The extra columns are the judge being stricter than the comparator prompt asks for — extra correct detail is supposed to be EQUAL. That is a calibration note, not a code change, and it is the kind of thing only reading the reasons will surface.

A failing case may mean the expectation is stale, not the model. Also on the first run, R8 (Kdo je prezident USA?) was marked WRONG for not calling `search_corpus`. The July prompt made every question route through retrieval; the current RULE 0 has a carve-out that makes general-knowledge questions tool-forbidden. So the model did what today's prompt says, and the case encodes yesterday's. **Decide which behaviour you want, then change one of them** — do not quietly relax the case, because a suite that is edited to stay green is worse than no suite.

What is not covered yet

- **Image cases.** V1-V7 and D3/D4 need image input; the collector's text path cannot send an attachment, so they are skipped with a note rather than silently passed.
- **Multi-turn.** Every case is a fresh conversation. Context-dependent drift — the kind that appears only after a model switch mid-chat — is not tested.
- **Statistical power.** One run per case. Sampling is non-greedy, so a single DRIFTING may be a draw rather than a trend. Re-run before believing it.

See also

- `docs/prompts.md` — the library, and how to add a case
- `tests/promptlib.py --validate` — the linter
- `scripts/agent.py --help` — the harness both phases drive
- `docs/measurements/` — the raw numbers behind the August prompt fixes

Chapter 9 Ingestion pipeline

Ingestion pipeline — mixed docx / xlsx / pdf corpus

Author: David Kubicek (david.kubicek@eywo.cz)

This is the design reference for how `services/ingester/ingester.py` (+ `xlsx_tables.py`, `corpus_manifest.py`) turns the corpus into Qdrant points. Those files are **frozen and regression-locked** (see BUILD_STATE step 04): treat this doc as the explanation of what the parser does and why, and the **code as the source of truth where they differ**. Do not edit the parser to match this doc — raise a blocker instead. The corpus is 10–30 files, mostly Czech + English, split between table-structured xlsx/csv (host → IP → VLAN → KVM-host assignments and contact sheets) and prose docx/pdf (descriptions in both languages).

The guiding principle: **tables and prose have different information shapes and are chunked differently**. A table row is a self-contained record; a paragraph is part of a flowing argument. One chunker for both would wreck retrieval on at least one.

Relevant ADRs: **ADR-0002** (BGE-M3 dense+sparse embedder, CPU), **ADR-0006** (content-addressed corpus sync — manifest, alias, watcher), **ADR-0007** (section-aware prose chunking; the `text/text_clean` split).

1. File routing

Walk the corpus root recursively (`scan_disk`). Route by extension (`iter_file`):

- `.xlsx`, `.xls` → **table path**, delegated to `xlsx_tables.iter_xlsx` (§2)
- `.csv` → **table path**, single-header CSV (`iter_csv`, §2)
- `.docx` → **docx path**: embedded tables → table path, body → prose path (§3)
- `.txt` → **prose path**, plain (`markdown=False`)
- `.md`, `.markdown` → **prose path**, Markdown headings honoured (`markdown=True`)
- `.pdf` → **prose path**, text-layer extraction (`pdfplumber`); a scanned PDF with no text layer is skipped with a warning
- anything else → skip with a logged warning (do not guess)

2. Table path (xlsx / csv / docx-tables) — one row per chunk

Tables are lookup data. The retrieval need is "give me the row for host X", "which hosts are on VLAN 203", or "the contact for the project lead" — precise, record-level. So:

- **One Qdrant point per data row**. Never embed a whole sheet or table as one chunk — that destroys row-level precision and overflows context.

- Each row is serialized to a flat, self-describing string joining each cell to its **header key** (`_serialize_kv`), sheet-prefixed for multi-sheet/table sources:

sheet: CutOver kontakty, Řízení projektu Jméno: Marek Česal, Řízení projektu
Firma: EPC, Tel: 739223474, E-mail: marek.cesal@epcommodities.cz

- Empty cells skipped; whitespace normalised. **Values kept verbatim** — no lowercasing/reformatting of IPs/hostnames/VLAN IDs; exact tokens are the point.
- This serialization is what makes the sparse vector useful: literal tokens (`prod-kvm-03`, `10.20.1.43`, `203`) land in the lexical index for exact match.
- **Payload** (per row): `source_file`, `sheet` (sheet name, "csv", or "table<n>" for docx tables), `row_index`, `kind`: "table_row", `lang`: "n/a" (rows are language-neutral key:value), and the serialized string as both `text` and `text_clean` (rows have no decoration to strip).
- **CSV** (`iter_csv`): first non-empty row is the header; a sheet with no sensible header is logged and skipped rather than guessed.

Messy real-world XLSX

All non-trivial XLSX structure — multi-table sheets, stacked/banded headers, multi-row title+header blocks, carry-forward and vertical-merge grouping, headerless data, trailing totals/notes trimming — lives in `xlsx_tables.py` and is locked by the offline regression (`tests/fixtures`, `test_xlsx_tables.py` → ALL PASS). This doc does not restate that logic; the code and its fixtures are authoritative. BGE-M3 handles 8192 tokens so a wide row is not a truncation risk; the parser front-loads identifying columns so the lexical signal leads.

3. Prose path (docx / pdf / txt / md) — section-aware chunks

(ADR-0007) Descriptions are flowing text where meaning spans sentences, and the primary docx uses **bold body text as headings** (not real Heading styles). The earlier paragraph-splitting chunker collapsed such a document into one page-sized slab whose embedding was a semantic average; ADR-0007 replaced it with section-aware chunking.

- **Section detection** (`_sections_from_lines`, `_docx_heading_level`): a new section starts at a Markdown `#...#####` heading, a Word Heading N/Title style, **or** a short standalone fully-bold docx paragraph (≤ 12 words, ≤ 80 chars — the bold-run fallback). Top-level heading opens a section and clears the subsection; a deeper heading becomes the subsection.
- **Keep a coherent section whole; split only when forced** (`_emit_sections`). Sizes are measured in **whitespace tokens (words)**, a deterministic proxy for model tokens:

- `CHUNK_MAX_WORDS = 480` (600 model tokens) — a section at or under this is emitted as **one chunk**, never split just to hit a target (no partial answers).
- `CHUNK_TARGET_WORDS = 300` — packing target when a larger section is split.
- Splitting happens at **sentence boundaries only** (`_SENT_SPLIT`); a chunk never cuts a sentence, so retrieval can't hand the model a truncated instruction. A lone sentence over the ceiling is emitted whole.
- `OVERLAP_FRAC = 0.15` — trailing whole sentences (15% of target) are carried into the next chunk so a topic straddling a split appears in both.
- **Two texts per chunk** (ADR-0007): `text` / `text_raw` is **verbatim** (returned to the LLM for citation/quoting); `text_clean` is Markdown-decoration-stripped (`_strip_md`) and is the **embedding input** so syntax stops polluting the dense vector. Link/URL tokens and hostnames are kept in `text_clean` (the sparse branch needs them); only scaffolding (`**`, `#`, list/quote markers, code fences, pipes) is removed. For non-Markdown sources the two coincide.
- **Line-span citation metadata**: `chunk_start_line` / `chunk_end_line`. docx has no source lines, so the **paragraph ordinal** (1-based over all paragraphs) is the line proxy; pdf uses extracted-text line numbers.
- **Language** (`detect_lang` via `langdetect` ON `text_clean`): `cs` | `en` | `other` | `unknown`. Metadata only — BGE-M3 is multilingual so retrieval does not depend on it, but it aids debugging/optional filtering.
- **Payload** (per chunk): `source_file`, `section_title`, `subsection_title`, `heading` (= `section_title`, used as retrieval's location), `chunk_index`, `chunk_start_line`, `chunk_end_line`, `kind`: "doc_text", `lang`, `text`, `text_clean`.

4. Embedding + storage

For each chunk (either path), batched (`BATCH = 64`):

1. Call the embedder `POST /embed` with `kind="passage"` on the chunk's `text_clean` (table rows fall back to their verbatim text).
2. Receive dense (1024-dim) and sparse (token→weight map) per chunk.
3. Upsert one Qdrant point with BOTH named vectors and the §2/§3 payload:
 - dense: 1024-dim, cosine distance
 - sparse: into Qdrant's sparse index
 - the exact embedded string is mirrored back into the payload as `text_clean`.

Point IDs are opaque `uuid4`. The parser also computes a deterministic per-row hash, but it is ignored on upsert — identity and pruning live in the checksum manifest (§5, ADR-0006), not the point ID. Old and new versions of a file never share an ID, which is what makes blind delete-by-recorded-UUID safe.

Collection schema + alias (created on first run)

```
vectors:
  dense: { size: 1024, distance: Cosine }
sparse_vectors:
  sparse: {}                                # Qdrant sparse index
```

Physical collections are named `corpus_<timestamp>_<uuid8>`; **retrieval targets the alias `corpus`** (QDRANT_COLLECTION), never a physical name. `--recreate` builds a new physical collection alongside the live one and atomically repoints the alias.

5. Sync, idempotency & re-runs (ADR-0006 — content-addressed)

Identity is the **file content checksum** (blake2b), tracked in a SQLite manifest (`corpus_manifest.py`, `manifest.db`) mapping checksum → {uuids, path, departed_at}. Identical-content files collapse to one checksum (intended dedup). Three modes (`main`, mutually exclusive, each under a pass lock so a manual run and the watcher never collide):

- **default — incremental** (`incremental`, `grace=None`): in-place sync of the live collection against disk. Order is **embed-before-prune** so no file's data is ever missing mid-pass: (1) add new checksums, (2) reconcile still-present content — cancel stale departures, and on a pure rename just refresh `source_file` in the payload (no re-embed), (3) prune departed checksums immediately.
- **--recreate** (`recreate`): full rebuild into a fresh `corpus_<ts>_<uuid8>`, embed everything, then **atomically switch the alias** and drop the old physical collection — **zero retrieval blackout**. The one-time exception is migrating a legacy pre-ADR-0006 physical `corpus` collection (a single momentary switch). Use `--recreate` whenever the schema, embedder model, or chunking changes.
- **--watch** (`watch`): autonomous debounced watcher (the `ragfarm-ingester-watcher` service). Debounce/quiescence (`INGEST_DEBOUNCE`) so a half-written file is never read and event storms coalesce; a full re-scan every `INGEST_SCAN_INTERVAL` as the missed-event backstop; **deletes are grace-gated** (`INGEST_DELETE_GRACE`) — a vanished file is only marked, and reappearing before the grace expires cancels the delete (rides out atomic-save / `rsync` windows). Read-only filesystem events (the watcher's own checksum reads) are ignored to avoid a self-feeding loop.

A model change (dim or model id) mixed into an existing collection is silently wrong; the alias/`--recreate` design makes the correct action (rebuild + switch) the easy one. Keep `models/embeddings/MODEL.md` in step with what a collection was built against.

6. Why hybrid (dense + sparse) for this corpus specifically

- **Dense** carries semantics and crosses languages: a Czech query finds the relevant English description and vice versa (the half the old English-only NPU build could never do). It embeds `text_clean`, so Markdown syntax doesn't skew it.
- **Sparse** carries exact lexical matches: querying `prod-kvm-03` or `10.20.1.43` hits the precise table row even though those tokens have no "semantics" to embed. Pure dense retrieval is unreliable for identifiers; sparse fixes it — which is why `text_clean` deliberately keeps URL/host tokens.
- `search_corpus` (rag-retrieval) fuses both with RRF over a broad candidate pool, then re-ranks with a GPU cross-encoder (`bge-reranker-v2-m3` via a dedicated `llama.cpp/Vulkan` server on `:8081`, ADR-0008), so one tool call serves both "which VLAN is host X on" (lexical) and "how do we handle host maintenance" (semantic, either language).

7. The automatic watcher (deployed autonomous sync)

This expands §5's `--watch` mode into the full operational picture — it is how the corpus stays in sync in normal running, with **no manual ingest**.

Deployment

The watcher runs as the systemd service `ragfarm-ingester-watcher` (host, as dave): `ingester.py --watch`, `Restart=on-failure`, logs to `logs/ingester-watcher.log`. It Wants/After `ragfarm-embedder.service` because every pass calls `/embed` — the embedder must be up first. It reads `CORPUS_PATH=/data/corpus`, writes through the alias `corpus` (`QDRANT_COLLECTION`), and keeps its manifest in `services/ingester/manifest.db` (`MANIFEST_DB`). Start/stop/status like any host service (see `docs/deployment.md` → Autostart).

How a pass is triggered

A background thread watches `CORPUS_PATH` recursively (watchdog **inotify** Observer; set `INGEST_WATCH_POLL=1` to use the `PollingObserver` on a network / overlay filesystem where `inotify` is unreliable). Two things schedule a pass:

- **Event + debounce.** A content-mutating event marks the state dirty a pass runs only after `INGEST_DEBOUNCE` seconds of **quiescence**, so a half-written file is never read and an event storm coalesces into one pass. **Read-only events are ignored** — every pass opens and reads each file to checksum it, which the observer sees as `opened/closed_no_write`; reacting to those would let the watcher's own hashing re-arm the debounce and spin forever (a self-feeding loop).

- **Full-scan backstop.** Regardless of events, a full re-scan runs every `INGEST_SCAN_INTERVAL` seconds, catching anything inotify missed.

What a pass does

Exactly the §5 `incremental()` sync, but with **grace-gated deletes** (`grace=INGEST_DELETE_GRACE`), embed-before-prune ordered so retrieval never sees a gap mid-pass:

1. **Add** new checksums (parse → embed → upsert, new uuid4 IDs, manifest updated).
2. **Reconcile** still-present content: cancel a pending departure if a file came back; on a pure rename just refresh `source_file` in the payload — **no re-embed**.
3. **Depart**: a vanished checksum is only marked with a timestamp, not deleted.
4. **Reap**: content absent for $\geq$ grace is deleted; reappearing before grace expires cancels the delete. This rides out atomic-save / rsync / editor-swap windows where a file briefly disappears and returns.

Concurrency & robustness

- Every pass takes a **non-blocking lock** on the manifest. If a manual run (`--recreate` or a manual incremental, which take the lock blocking) is in progress, the watcher **skips that tick and retries next** — the two never double-write. After a manual `--recreate` flips the alias, the watcher's next pass automatically operates on the new physical collection (it resolves the alias each pass), so the two modes compose cleanly.
- A failed pass is logged and the loop continues; a crashed process is restarted by systemd (`RestartSec=5`).
- The watcher only does **incremental** sync — it never `--recreates`. A schema, embedder-model, or chunking change still needs a manual `--recreate` (§5); the watcher then picks up the new collection via the alias.

Deployed knobs (unit values; all env-overridable)

env	unit value	meaning
<code>INGEST_DEBOUNCE</code>	60 s	quiescence required after an event before a pass
<code>INGEST_DELETE_GRACE</code>	120 s	how long a vanished file is tolerated before its points are pruned
<code>INGEST_SCAN_INTERVAL</code>	3600 s	full re-scan backstop for missed events

INGEST_WATCH_POLL	unset (inotify)	1 → polling observer for network/overlay FS
-------------------	-----------------	---

8. Verification (feeds step 04's gate)

After ingesting the corpus (or a 2–3 file subset):

- The offline parser regression passes: `FIXTURES=tests/fixtures python services/ingester/test_xlsx_tables.py` → ALL PASS.
- Alias corpus resolves to a physical collection with point count > 0 and a schema showing dense (1024) + sparse.
- A query for a known hostname returns its exact table row in the top hits (sparse path works).
- A Czech semantic query returns a relevant doc chunk (multilingual dense works). Together these prove the whole reason for the BGE-M3/CPU redo and the ADR-0006/0007 rework.

Chapter 10a Embedder README

Embedder — BAAI/bge-m3 on CPU (:8090/embed)

Per ADR-0002 the embedder runs **BAAI/bge-m3 on CPU** (the NPU path was abandoned — English-only, seq-limited models don't fit this mixed Czech/English, wide-table corpus; `quark_quantize.py` here is a vestige of that abandoned NPU attempt).

One endpoint, one job

This service exposes exactly **one** endpoint: `POST /embed`. It returns BGE-M3's dense (1024-dim) + sparse vectors in a single pass, for both ingestion and query embedding. That is all it does.

Reranking is NOT here. It used to be a `/rerank` sub-endpoint on this service, but the reranker and the embedder no longer share a model family, a device, or a purpose, so co-hosting them was a wart (ADR-0008). The cross-encoder reranker is now its own dedicated GPU service — a `llama-server --reranking` instance on **:8081** — with its own unit (`manifests/ragfarm-reranker.service`) and record (`models/reranker/MODEL.md`). This service stays embeddings-only.

Contract

```
POST http://127.0.0.1:8090/embed
  {"input": ["text1", ...], "kind": "passage|"query"}
->{"dense": [...1024...], "sparse": [{"<tok_id>": weight, ...}], "dim": 1024}
```

Run

Code: `services/embedder/server.py` (FastAPI + `FlagEmbedding BGEM3FlagModel`).
Unit: `manifests/ragfarm-embedder.service` (host, CPU; `MBED_MODEL_PATH` pins the BGE-M3 snapshot). Model record + revision: `models/embeddings/MODEL.md`.

Chapter 10b llama.cpp (reranker engine)

llama.cpp — CUDA build, for the reranker only

Scope, because this used to say something else. llama.cpp is no longer the generation engine. ADR-0013 moved generation to vLLM, and every LLM in this deployment is served by ragfarm-vllm@N on the OpenAI-compatible API. What survives on llama.cpp is exactly one service:

```
ragfarm-reranker — bge-reranker-v2-m3 ON llama-server --reranking, 127.0.0.1:8081/
reranking. A cross-encoder, not a generator. See ADR-0008.
```

Everything about AMD, Vulkan, iGPU offload, --mmaproj vision projectors and GGUF generation models has been removed from this file. It described the retired box and was actively misleading next to a CUDA deployment. Git history has it if you ever need the provenance.

Why keep llama.cpp at all: --reranking gives a small, well-behaved cross-encoder server with no Python runtime, negligible memory, and no competition for vLLM's memory budget. Replacing it with a second vLLM instance would cost more GPU than the reranker itself uses.

Build

CUDA, not Vulkan. The target is a GB10 Grace Blackwell, compute capability **sm_121**.

```
git clone https://github.com/ggml-org/llama.cpp "$LLAMA_DIR"
cmake -S "$LLAMA_DIR" -B "$LLAMA_DIR/build" -DGGML_CUDA=ON -
DCMAKE_BUILD_TYPE=Release
cmake --build "$LLAMA_DIR/build" --config Release -j "${MAX_JOBS:-8}"
```

Two things that bite:

- **MAX_JOBS is load-bearing.** Unbounded, ninja runs nproc+2 parallel CUDA compiles and the OOM killer takes the build at 98 GB. Same lesson as the vLLM build; a different budget from --gpu-memory-utilization entirely.
- **\$LLAMA_DIR** is where the rest of the tooling expects to find it. scripts/deploy.sh and scripts/fetch-encoder.sh both check for \$LLAMA_DIR/build/bin/llama-server and refuse to continue without it. The reranker unit currently hardcodes its path rather than reading LLAMA_DIR — a known inconsistency, noted in PROGRESS.md.

Verify:

```
"$LLAMA_DIR/build/bin/llama-server" --version
scripts/stack.sh status           # reranker row should read [OK]
```

Running it

You should not need to. `ragfarm-reranker.service` owns it, and `scripts/stack.sh` starts, stops and health-checks it with everything else. The unit runs at a mildly lowered scheduling priority so the interactive LLM wins any contention — the reranker is the most expensive stage of retrieval (297 ms) but it is not the one a human is waiting on keystroke by keystroke.

By hand, for debugging only:

```
"$LLAMA_DIR/build/bin/llama-server" --reranking \  
-m "$RERANK_MODEL_PATH" --host 127.0.0.1 --port 8081  
curl -s http://127.0.0.1:8081/health
```

See also

- `models/reranker/MODEL.md` — which checkpoint, and why full quality
- `docs/decisions/ADR-0008` — why a cross-encoder reranker at all
- `docs/decisions/ADR-0013` — the engine split that retired `llama.cpp` for generation
- `man docs/man1/stack.1` — lifecycle and health checking

Chapter 11 Configuration reference

(.env.example)

Reproducible source-of-truth for every environment variable in the ragfarm stack. Copy to .env on the host and set the ones you need — everything else stays commented at safe defaults.

```
# Repo-root .env – copy to .env and fill in. The real .env is gitignored and
# lives only on the build host. Plain VAR=VALUE lines; no quoting needed.
#
# -----
# NEW HERE? READ THE MAN PAGE FIRST:   man docs/man1/env.1
#
# It covers what this file cannot: the four DIFFERENT mechanisms that carry a
# value to its consumer, which one applies to what, the precedence order, the
# two virtualenvs (.venv vs .venv-vllm – there is no .venv0/.venv1), and the
# procedure for changing a setting so that it actually takes effect. THIS file
# stays the single source of truth for what each individual variable MEANS.
#
# Related pages:  man docs/man1/activate_llm.1      slots and the GPU budget
#                 man docs/man1/fetch_llm.1        getting models onto the box
#                 man docs/man1/active.json.1       the model registry
#
# CHANGED RECENTLY – check these if you are returning to the project:
#   2026-08-09  LLM_REASONING_PARSER is RETIRED as a global. It is now
#               LLM_REASONING_ARGS, written per slot into .env.slotN by
#               activate_llm.py from the registry's `profile` field. A global
#               value applied the Thinking parser to an Instruct slot and the
#               entire answer came back filed as "reasoning", leaving the chat
#               message empty. See the note at that variable.
#   2026-08-09  RAGFARM_PUBLIC_HOST added. The ONE address in this system that
#               must not be 127.0.0.1, because a browser resolves it. Two
#               consumers must agree on it; compose cannot autodetect it.
#   2026-08-09  max_tokens for both vision profiles raised to 24576 (that lives
#               in setup_openwebui.py, not here – noted so you know where).
#
# -----
#
# HOW THIS FILE REACHES THE RUNNING SYSTEM:
#   `source scripts/proxy-env.sh` loads EVERY line here into the shell (not just
#   the proxy vars – the loader exports all of them). Any var referenced in
#   infra/compose.yaml as `${VAR:-default}` then flows into containers via that
#   shell env when you run `docker compose up`. Always source the loader first:
#       source scripts/proxy-env.sh && docker compose -f infra/compose.yaml up -d
#   Vars NOT listed as `${VAR}` in compose.yaml (Category 2 below) are baked in
#   as literals there and are NOT affected by anything you put in this file when
#   running via compose – they only apply if you run a service's server.py
#   directly on the host. This distinction is exactly what caused the confusion
#   this file exists to resolve – read the category headers below.
#
```

```

# Full variable inventory generated by grepping every os.environ/getenv,
# `${VAR}` compose substitution, and shell `${VAR}` read in the repo.

# =====
# 0. CANONICAL ENDPOINT NAMES (added 2026-08-03)
# .env AT THE REPO ROOT IS THE SINGLE SOURCE OF TRUTH. Python code should read
# these through `ragfarm_env.py` (repo root) rather than calling os.environ
# with a hand-typed name – that is what keeps the naming consistent.
# python ragfarm_env.py # print everything as resolved, incl. warnings
#
# RULE: base URLs (scheme://host:port, NO path) are the source of truth. Full
# endpoints are derived in ragfarm_env.py (EMBED_ENDPOINT = EMBED_URL+/embed).
#
# Reaching consumers:
# host services -> systemd `EnvironmentFile=-/home/dave/ragfarm/.env`
# containers -> `env_file: ../.env` in infra/compose.yaml
# (compose's own .env handling only does `${VAR}`
# substitution
# in the compose file; it never puts the file in an
# image)
# scripts/tools -> `import ragfarm_env` (calls load_dotenv itself)
#
# DEPRECATED SPELLINGS, still honoured, reported by
# `ragfarm_env.deprecations()`:
# LLAMA_URL -> LLM_URL LLM_GGUF_PATH -> LLM_MODEL_PATH
# RERANK_GGUF_PATH -> RERANK_MODEL_PATH LLM_GGUF_MMPROJ -> LLM_MMPROJ_PATH
# The *_GGUF_* names are format-specific and wrong once vLLM serves
# safetensors
# (ADR-0013); LLM_MODEL_PATH is format-neutral (a GGUF file OR a snapshot
# dir).
# =====
#LLM_URL=http://127.0.0.1:8080 # generative LLM, OpenAI-compatible
#RERANK_URL=http://127.0.0.1:8081 # cross-encoder reranker (ADR-0008)
#EMBED_URL=http://127.0.0.1:8090 # BGE-M3 dense+sparse embedder
#QDRANT_URL=http://127.0.0.1:6333
#MCP0_URL=http://127.0.0.1:8000 # MCP->OpenAPI bridge; tools mount at /
#<name>
#RAG_MCP_URL=http://127.0.0.1:8104 # rag-retrieval MCP (behind mcp0;
# use :8000/rag)
#OWUI_URL=http://127.0.0.1:3000
#QDRANT_COLLECTION=corpus

# Materialized full endpoints. Normally DERIVED from the bases above, but spelled
# out because two consumer classes cannot import ragfarm_env: the FROZEN ingester
# (reads EMBED_ENDPOINT directly) and the containers. KEEP CONSISTENT with the
# bases – `python ragfarm_env.py` warns via drifts() if they diverge.
#EMBED_ENDPOINT=http://127.0.0.1:8090/embed
#RERANK_ENDPOINT=http://127.0.0.1:8081/reranking

```

```

# Ingester watch tuning (moved out of ragfarm-ingester-watcher.service 2026-08-03
-
# units carry start/stop/restart POLICY only, never config).
#INGEST_DELETE_GRACE=120
#INGEST_SCAN_INTERVAL=3600
#INGEST_DEBOUNCE=60

# =====
# 1. LIVE KNOBS – vars that actually change running behavior. Two mechanisms:
#   1a. containers: infra/compose.yaml `${VAR:-default}` substitution.
#   1b. host systemd units (llama/embedder/reranker/ingester-watcher – NOT
#       compose): each unit has EnvironmentFile=~/home/dave/ragfarm/.env and, as
#       of 2026-08-03, NO Environment= lines at all – units carry start/stop/
#       restart POLICY only, config lives here. A unit whose value is missing
#       here now fails loudly instead of silently using a stale baked-in default
#       (the llama unit used to pin a Qwen2.5 GGUF path that way). Restart to
#       apply.
# =====

# --- 1a. corpus (step 04 ingester) -----
# Absolute path to the read-only corpus on the host. Also read directly by
# services/ingester/ingester.py when run bare (BUILD_STATE step 04 passes
# --corpus explicitly instead, which overrides this – see BUILD_STATE.md).
CORPUS_PATH=/data/corpus

# --- OpenNebula MCPs: mock mode + safety allowlist (ADR-0004) -----
# ONE MOCK / HOST MOCK: "1" (default) serves canned data / simulates actions
# with NO live OpenNebula – lets the agent path and confirmation UX be tested
# with no cluster. Set to "0" only once ONE_XMLRPC/ONE_AUTH are real (see
# services/mcp-placement/.env.example) and HOST_ALLOWLIST is deployment-real.
#ONE MOCK=1
#HOST MOCK=1
# HOST_ALLOWLIST: comma-separated hostnames mcp-host-control's reboot_host will
# ever act on with confirm=True – a target-hostname allowlist, NOT a user
# allowlist (there is currently no per-user gate; see ADR-0004 §4 / the
# confirmation-wrapper pattern for who can trigger it). Refuses any host not
# listed, even if the wrapper's confirmation modal was approved.
#HOST_ALLOWLIST=node-01,node-02,node-03

# --- Open WebUI exposure -----
# 0.0.0.0 (default) = reachable from the LAN, login-gated by WEBUI_AUTH – the
# only service meant to be externally reachable. Set 127.0.0.1 for loopback-only
# (reach via `ssh -L 3000:127.0.0.1:3000 <host>` instead). Restrict :3000 at the
# host firewall to trusted networks regardless of this setting.
#OWUI_HOST=0.0.0.0

# --- Public host: the address a REMOTE BROWSER uses to reach this box -----
# Everything else in this repo is server-to-server and correctly says 127.0.0.1.

```

```

# This one is different: the draw.io wrapper in the vision preset's RULE 5 is
# fetched by the USER'S browser, inside OWUI's HTML-preview iframe. On a remote
# client 127.0.0.1 is the client's own loopback, so viewer-static.min.js never
# loads and the preview pane renders an empty white box – with no error anywhere.
#
# Two consumers must agree on this value:
#   - infra/openwebui/setup_openwebui.py bakes it into the RULE 5 wrapper URLs.
#     If unset it AUTODETECTS from the default route and prints a warning.
#   - infra/compose.yaml puts it in IFRAME_CSP so the iframe is allowed to load
#     from it. Compose CANNOT autodetect: if this is unset it falls back to
#     127.0.0.1 and the CSP silently blocks the (correctly autodetected) URL.
# So: set it here. Autodetect is a safety net for a loopback-only box, not the
# supported configuration. Changing it needs `docker compose up -d open-webui`
# (CSP is read at container start) AND a setup_openwebui.py re-run.
#
# Use whatever the audience's browser can actually resolve – the LAN address of
# this box normally, but a VPN/NAT address or a DNS name if that is how they
# reach it. Hostname is fine; it is used verbatim in a URL.
#RAGFARM_PUBLIC_HOST=192.168.0.208

# --- Prompt-round timeout ceilings (all seconds) -----
# OWUI's aiohttp client uses AIOHTTP_CLIENT_TIMEOUT for /v1/chat/completions and
# outbound tool-server calls; its built-in default is 300.
AIOHTTP_CLIENT_TIMEOUT_TOOL_SERVER
# defaults to AIOHTTP_CLIENT_TIMEOUT unless set. mcpo's CONNECTION_TIMEOUT is
# used as
# SSE-read timeout for MCP-subprocess tool calls. rag-retrieval's embedder and
# reranker
# HTTP calls are hardcoded at 300 s (services/rag-retrieval/server.py). llama-
# server's
# --timeout is 3600 s by default (already ample).
#
# Bumped 600 -> 1200 s on 2026-07-27: Qwen3-VL-Thinking OCR + translation prompts
# were pushing 530 s wall-clock; 1200 s gives comfortable headroom for the
# current
# iGPU throughput and stays well below llama-server's own 3600 s ceiling.
# Set in infra/compose.yaml; values here override on the shell for one-off runs.
#AIOHTTP_CLIENT_TIMEOUT=1200
#AIOHTTP_CLIENT_TIMEOUT_TOOL_SERVER=1200
#CONNECTION_TIMEOUT=1200

# --- outbound proxy (optional) -----
# Set these when the build host reaches PyPI / HuggingFace / container
# registries through a corporate proxy. Leave unset/blank for direct access.
#HTTP_PROXY=http://proxy.corp.example:3128
#HTTPS_PROXY=http://proxy.corp.example:3128

# Hosts that must BYPASS the proxy. The loader always prepends the internal
# baseline (localhost,127.0.0.1,0.0.0.0,::1,host.docker.internal,qdrant), so

```

```

# list here only ADDITIONAL no-proxy targets – e.g. your LAN / OpenNebula subnet.
#NO_PROXY=.corp.example,10.0.0.0/8

# --- 1b. host model units: LLM (ragfarm-vllm), embedder (ragfarm-embedder),
#     reranker (ragfarm-reranker) – see the 1b note above for the override rule
# ---
#
# MAX_JOBS: FIRST-RUN LIMITER for JIT kernel builds. Not a steady-state tuning
# knob – set it for a cold box, then unset it once the caches are warm.
#
# Why it exists: FlashInfer JIT-compiles CUDA kernels at RUNTIME on first use,
# and
# flashinfer/jit/cpp_ext.py only passes `-j` to ninja when MAX_JOBS is set.
# Unset,
# ninja defaults to nproc+2 (22 here) parallel CUDA compiles; CUTLASS translation
# units are multi-GB each, so on the Spark's SHARED 128 GB that peaked at ~98 GB
# and the OOM killer took ragfarm-vllm three times. It reads as a slow model
# load,
# because it dies AFTER all four shards are in.
#
# NOTE this is a different budget from --gpu-memory-utilization
# (LLM_GPU_MEM_UTIL):
# that one bounds weights+KV at STEADY state; MAX_JOBS bounds COLD-START
# compilers.
# Lowering the vLLM memory flags does nothing for this peak – the memory is the
# compilers', not vLLM's.
#
# When can you unset it? Once both caches are populated, startup skips
# compilation
# and MAX_JOBS stops mattering (warm restart measured 3m11s vs 813s cold):
#   ls ~/.cache/flashinfer/           # expect a version dir, e.g. 0.6.14/
#   du -sh ~/.cache/flashinfer/       # expect hundreds of MB once built
#   ls ~/.cache/vllm/torch_compile_cache/ # torch.compile/inductor graphs
# A warm start logs "Directly load the compiled graph(s) ... from the cache".
# Leaving it set is harmless – it only slows a rebuild that rarely happens. Re-
# set
# it before any upgrade of vllm/flashinfer, a CUDA bump, or a new checkpoint,
# since
# each invalidates the JIT cache and you are cold again.
MAX_JOBS=4
#
# LLM_GGUF_PATH: main GGUF. Written by `scripts/fetch-llm.sh` (download) or
# `scripts/activate-llm.sh` (switch among models already on disk).
# NOTE (ADR-0013): GGUF is the retired llama.cpp GENERATION path. vLLM serves
# NVFP4 safetensors via LLM_MODEL_PATH; see the vLLM block further down.
#LLM_GGUF_PATH=/home/dave/ragfarm/models/llm/qwen2.5-7b-instruct-gguf/qwen2.5-7b-
instruct-q4_k_m-00001-of-00002.gguf
# LLM_GGUF_MMPROJ: multimodal-projector GGUF, only for a vision-language model
# (e.g. Qwen2.5-VL) – blank for text-only. fetch-llm.sh/activate-llm.sh detect

```

```

and
# set this automatically (from a repo's own *mmproj*.gguf); no manual flag
needed.
# They also CLEAR it when the active model has none, so a stale path never
lingers.
#LLM_GGUF_MMPROJ=
#
# --- vLLM serving (ADR-0013, step 02). Read by manifests/ragfarm-vllm.service
---
# LLM_MODEL_PATH: NVFP4 safetensors SNAPSHOT DIR (not a file).
#LLM_MODEL_PATH=/home/dave/ragfarm/models/llm/Qwen3-VL-30B-A3B-Instruct-NVFP4
# LLM_SERVED_NAME: LOAD-BEARING. infra/openwebui/setup_openwebui.py keys
# MODEL_TUNING by whatever alias /v1/models advertises; serve under another name
# and the ragfarm-vision preset silently fails to bind.
#LLM_SERVED_NAME=qwen3-vl-30b-a3b
# LLM_MOE_BACKEND: NVFP4 MoE kernel. `auto` is recommended – it re-evaluates what
# is actually available, which changes with the JIT cache state. Do NOT use plain
# `flashinfer`: not a valid value (argparse rejects it). Real values include
# flashinfer_b12x / flashinfer_cutlass / cutlass / marlin. `marlin` is the LAST
# resort – it dequantizes FP4->BF16 and forfeits the FP4 speedup.
#LLM_MOE_BACKEND=auto
#LLM_MAX_MODEL_LEN=32768
# LLM_REASONING_ARGS: do NOT set this here. It is written per slot into
# .env.slotN by scripts/activate_llm.py, from the registry entry's profile,
# because the reasoning parser is a property of the CHECKPOINT rather than of
# this box – Thinking models need `--reasoning-parser qwen3`, Instruct models
# must have no parser at all.
#
# There used to be a global LLM_REASONING_PARSER=qwen3 here, and on 2026-08-09 an
# Instruct slot inherited it. The qwen3 parser splits the stream on </think>, and
# with no </think> in it – an Instruct model never emits one – it classifies the
# ENTIRE answer as reasoning and leaves content empty. The chat message comes
back
# blank with the whole diagram hidden inside the collapsible "thinking" trace,
# which is convincingly mistakable for a broken model or a bad quantisation.
# A global value cannot express a per-checkpoint fact; that is the whole lesson.
# LLM_GPU_MEM_UTIL: steady-state weights+KV budget. MUST be set explicitly here:
# unified memory has no private VRAM, so vLLM's 0.9 default reserves ~108 GB of
the
# shared 128 GB and starves the embedder, reranker, Qdrant and the OS.
# At 0.50 the measured KV cache is 36.52 GiB (398,832 tokens).
#LLM_GPU_MEM_UTIL=0.50
# Batch caps. Bound the startup profile_run and steady-state activations; on a
30B
# MoE the profile pass materialises per-expert workspaces.
#LLM_MAX_NUM_BATCHED_TOKENS=8192
#LLM_MAX_NUM_SEQS=64
#
# EMBED_MODEL_PATH: embedder weights dir. Written by `scripts/fetch-encoder.sh`.

```

```

# Swapping this changes the vector space – re-ingest is required (see
deployment.md
# → "Activating a swapped model").
# EMBED_MODEL_PATH=/home/dave/ragfarm/models/embeddings/bge-m3
# RERANK_MODEL_PATH: reranker weights, also written by `scripts/fetch-encoder.sh`
# (fetches the matched embedder+reranker pair together). Still a GGUF – the
reranker
# is the one thing llama.cpp still serves (ADR-0013 §3) – but the key is now
# format-neutral like the others. The legacy RERANK_GGUF_PATH spelling is still
# honoured by ragfarm_env.py and reported by deprecations().
# RERANK_MODEL_PATH=/home/dave/ragfarm/models/reranker/bge-reranker-v2-m3/bge-
reranker-v2-m3-f16.gguf

# =====
# 2. INTERNAL WIRING – read by each service's code (os.environ.get with a
# default), but infra/compose.yaml sets these to a FIXED LITERAL per
# service, not `${VAR}` substitution. Putting them in .env has NO EFFECT on
# the containerized services. Listed here only so "what does this var do"
# has one answer – they matter if you run a server.py directly on the host
# (outside compose) for local debugging.
# =====
# QDRANT_URL=http://localhost:6333          # ingestor.py, rag-retrieval/server.py
# EMBED_ENDPOINT=http://localhost:8090/embed # ingestor.py, rag-retrieval/
server.py
# RERANK_ENDPOINT=http://localhost:8081/reranking # rag-retrieval -> dedicated GPU
llama.cpp reranker (ADR-0008)
# QDRANT_COLLECTION=corpus                  # ingestor.py, rag-retrieval/
server.py
# RAG_HOST=0.0.0.0                          # rag-retrieval/server.py bind
address
# RAG_PORT=8104                            # rag-retrieval/server.py bind port
# FS_SANDBOX_ROOT=/data/corpus             # mcp-fs/server.py (experimental,
unbridged)
# RAG_PREFETCH: candidates pulled per branch (dense/sparse) before RRF fusion in
# search_corpus. Not set anywhere in compose – code default (40) always applies
# unless exported manually before running rag-retrieval/server.py directly.
# RAG_PREFETCH=40
# --- rerank + window (ADR-0007/0008; rag-retrieval/server.py, code defaults
shown) ---
# All are unset in compose, so the code defaults below always apply unless
# exported manually. Tune these when evaluating retrieval quality.
# RAG_CANDIDATES: size of the fused RRF pool handed to the re-ranker (broad-in).
# Wide enough that all same-template rows (e.g. every EPC contact) are in play.
# RAG_CANDIDATES=40
# RAG_USE_RERANKER: 1 (default) = cross-encoder rerank via RERANK_ENDPOINT; 0 =
# legacy MMR path (kept only for A/B – see ADR-0008 for why MMR mis-fires on the
# small row-per-record chunks by treating N distinct records as "redundant").
# RAG_USE_RERANKER=1
# RAG_MIN_SCORE: absolute reranker-score floor (ADR-0010 §1 workhorse). Drop

```

```

# reranked hits whose normalized (sigmoid) relevance is below this. 0.0 = keep
all
# (default until calibrated on real dumps via `scripts/rag_pool_inspect.py
# --dump-scored calib.csv <queries...>` – label required/junk in the CSV and read
# the boundary). Observed on the corpus: real answers 0.13–0.95, reverse FW rows
# as low as 0.56, junk ~0.002 – expected calibrated value ~0.35–0.45.
#RAG_MIN_SCORE=0.0
# RAG_GATE_KNEEDLE: ADR-0010 $1 Kneedle adaptive hatch, applied AFTER the floor.
# 1 (default) = when survivors exceed RAG_GATE_MIN_SET and a strong chord-
distance
# knee exists in the sorted-score curve, cut at the elbow; degrades to a no-op on
# flat curves. Independent of RAG_LIGHTRAG (traversal branch, not built yet).
#RAG_GATE_KNEEDLE=1
# RAG_GATE_MIN_SET: arm the Kneedle hatch only when more than this many hits
# survive the floor – smaller pools do not need extra trimming.
#RAG_GATE_MIN_SET=12
# RAG_GATE_WEAK_KNEE: normalized chord distance below which the strongest knee is
# treated as "no clear elbow" and the set is left alone (fraction of score span).
#RAG_GATE_WEAK_KNEE=0.05
# RAG_MMR_LAMBDA: LEGACY, only used when RAG_USE_RERANKER=0. MMR relevance-vs-
# diversity weight in [0,1]; 0.3 leans toward diversity. Superseded by the
reranker.
#RAG_MMR_LAMBDA=0.3
# RAG_EXPAND_NEIGHBORS: same-section neighbor chunks to fold into each returned
hit
# on EACH side (0 disables window expansion). Only widens split sections; never
# crosses a section boundary.
#RAG_EXPAND_NEIGHBORS=1
# RAG_EXPAND_MAX_WORDS: word cap on an expanded window; neighbors past the cap
are
# dropped so a hit never balloons back into a page-sized slab.
#RAG_EXPAND_MAX_WORDS=600
# EMBED_MODEL_PATH and RERANK_MODEL_PATH are host-unit knobs, listed under 1b
# above (NOT here – they are not compose-literal like the rest of this category).

# =====
# 3. ONE-SHOT ADMIN SCRIPT OVERRIDES – infra/openwebui/setup_openwebui.py and
#   check_toolchain.py. Normally passed inline on the command line (they're
#   invoked ad hoc, not run as services), e.g.:
#   OWUI_EMAIL=admin@ragfarm.local OWUI_PASSWORD=... python3 infra/openwebui/
setup_openwebui.py
#   Can be placed here instead for convenience since the loader exports
#   everything – but OWUI_PASSWORD is a credential, so weigh that before
#   putting it in a file, even a gitignored one.
# ---- setup_openwebui.py ----
#OWUI_URL=http://127.0.0.1:3000
#OWUI_TOKEN=                                # admin JWT, alternative to
email+password
#OWUI_EMAIL=admin@ragfarm.local

```



```

#OWUI_PASSWORD=
#MCPO_RAG_URL=http://127.0.0.1:8000/rag
#MCPO_PLACEMENT_URL=http://127.0.0.1:8000/placement
#BASE_MODEL_ID=qwen2.5-7b-instruct
# ---- check_toolchain.py (step-07 gate check) ----
#MCPO_URL=http://127.0.0.1:8000
#LLAMA_URL=http://127.0.0.1:8080
#MCPO_CONTAINER=infra-mcpo
#OWUI_MODEL=ragfarm
#RAG_TOOL_ID=server:0
#HOSTNAME_PROBE=hsmbvxip001ts
#CZECH_PROBE=Jak se přistupuje z ŠA do hostingu?

# =====
# 4. TEST-ONLY
# =====
# FIXTURES: fixture dir for services/ingester/test_xlsx_tables.py (offline
# parser regression, no services needed). Default tests/fixtures is correct for
# a normal checkout; override only if running fixtures from elsewhere.
#FIXTURES=tests/fixtures

# =====
# Per-service secrets live in their own directory, not here:
#   services/mcp-placement/.env.example -> ONE_XMLRPC, ONE_AUTH (OpenNebula
#   creds)
# mcp-host-control currently has no secrets of its own (HOST MOCK/HOST_ALLOWLIST
# come from Category 1 above via compose), so it has no .env.example – per
# ADR-0005 collocation rules, one is added only if/when it takes secrets.
# =====

```

Chapter 12a Model record — LLM

Generative LLM Model Record — the Qwen3-VL fleet

Supersedes the Qwen2.5-7B-Instruct / GGUF-Q4_K_M / llama.cpp-Vulkan record that lived here (ADR-0001, AMD box). Generation moved to vLLM on the Spark per ADR-0013; llama.cpp still serves the reranker (../reranker/MODEL.md), and the embedder is in ../embeddings/MODEL.md.

There is no longer a single LLM. `active.json` in this directory is the git-tracked registry of every checkpoint the deployment should hold and which of them are bound to a vLLM **slot** `scripts/fetch_llm.py` and `scripts/activate_llm.py` are the only supported way to change it (ADR-0013 §2a). Do not hand-edit `LLM_MODEL_PATH` / `LLM_SERVED_NAME` — they now live in the generated `.env.slotN` files.

model	quant	alias / preset	why it is here
Qwen3-VL-30B-A3B- Instruct	NVFP4	...-instruct-nvfp4 / ragfarm-vision-instruct	first model on the Spark; fast (76 tok/s), weaker at multi-column reasoning
Qwen3-VL-30B-A3B- Thinking	NVFP4	...-thinking-nvfp4 / ragfarm-vision	flagship. Same quant as the Instruct, so Instruct-vs-Thinking is a clean single-variable comparison
Qwen3-VL-30B-A3B- Thinking	FP8	...-thinking-fp8 / ragfarm-vision-fp8	official Qwen build; quant-fidelity reference against the NVFP4 one
Qwen3-VL- 32B -Thinking	NVFP4	...-32b-thinking-nvfp4 / ragfarm-vision-dense	DENSE , same family and quant as the MoE — isolates MoE-vs-dense
Qwen3-VL- 8B -Thinking	bf16	...-8b-thinking / ragfarm-vision-8b	the old AMD box's quality benchmark, at bf16 rather than its Q4_K_M

Sizes, revisions and sources are in `active.json`; `scripts/fetch_llm.py --verify` byte-checks every one of them against the Hub.

Baseline record — Qwen3-VL-30B-A3B-Instruct (NVFP4)

Everything measured below was measured on THIS checkpoint. Treat it as the baseline the others are compared against, not as "the model".

Field	Value
Repo	ig1/Qwen3-VL-30B-A3B-Instruct-NVFP4 (community NVFP4 + vision)
Revision	3c6162d5513d26f008628eebe9b4355559b4a305 (resolved 2026-08-03, ungated)
Served alias	qwen3-vl-30b-a3b at the time of measurement; the registry now assigns qwen3-vl-30b-a3b-instruct-nvfp4. The alias is load-bearing — it is what OWUI binds its preset to
Architecture	Qwen3VLMoeForConditionalGeneration, model_type: qwen3_vl_moe
MoE shape	128 experts/layer, top-8 routed, 48 layers (all MoE). 4.7M params per expert → 29B total, 3B active per token
Quantization	compressed-tensors, format nvfp4-pack-quantized, weights 4-bit float. Vision blocks are in the ignore list (kept at higher precision)
Weights	4 safetensors shards, 17.86 GiB on disk
Engine	vLLM 0.26.0 in a DEDICATED venv .venv-vllm (torch 2.11.0+cu130)
Unit	manifests/ragfarm-vllm@N.service (templated per slot) → 127.0.0.1:8080/v1 for slot 0
Context	--max-model-len 32768
Updated	2026-08-05

MoE backend on sm_121 — MEASURED, and it contradicts ADR-0013

ADR-0013 targets --moe-backend flashinfer as "the b12x path". On this box:

- **flashinfer is not a valid flag value** in vLLM 0.26.0. The choices are flashinfer_b12x, flashinfer_cudtsl, flashinfer_cutlass, flashinfer_trtllm, cutlass, marlin, triton, ... Passing flashinfer is rejected by argparse.
- **flashinfer_b12x DOES work on GB10** — verified 2026-08-03. All three of _supports_current_device()'s conditions pass (is_cuda, is_device_capability_family(120), has_flashinfer_b12x_moe()), it serves, and it

generates coherent text including Czech — not the !!!!! mode. It is however **opt-in only**: vLLM's oracle excludes it from auto-selection "until the upstream CUTLASS SM121 MMA op guard is resolved".

- **What we run by default:** `FLASHINFER_CUTLASS`, via `--moe-backend auto`, with `FlashInferCutlassNvFp4LinearKernel` for the linear NVFP4 GEMM. Native FP4, ranked ABOVE MARLIN — the marlin dequant fallback was never needed on this box.

Measured decode (256 tok, `ignore_eos`, single stream, 3 runs): **flashinfer_b12x 75.6 tok/s** vs `FLASHINFER_CUTLASS 71.1 tok/s`. b12x is 6% faster, but that is decode — which is bandwidth-bound and so the least favourable comparison for FP4. Prefill is unmeasured. Default stays `auto`: 6% does not justify overriding an explicit upstream "not safe to auto-select yet" on a production box.

Do not conclude a FlashInfer backend is unsupported without checking PATH. An earlier run here reported b12x as does not support current device cuda and a different run picked `VLLM_CUTLASS` — both were artifacts of `ninja` missing from the unit's PATH, since `has_flashinfer_b12x_moe()` probes by import. Backend selection also shifts with JIT cache state, so one log line is not a fact.

Startup memory — the two budgets are NOT the same lever

Cold start was OOM-killed three times at a near-identical **98.3-98.6 GiB** peak with 13.6 GiB of swap, and the peak did not move when `--gpu-memory-utilization`, `--max-num-batched-tokens` and `--max-num-seqs` were lowered.

- **Cold-start build memory** — `MAX_JOBS` (in `.env`). FlashInfer JIT-compiles kernels at runtime, and `flashinfer/jit/cpp_ext.py` only passes `-j` to `ninja` when `MAX_JOBS` is set. Unset, `ninja` defaults to `nproc+2` = **22 parallel CUDA compiles**. CUTLASS translation units are multi-GB each. That, not vLLM, was the 98 GiB. `MAX_JOBS=4` fixes it. Only bites on a cold JIT cache (`~/.cache/flashinfer/`); warm starts skip compilation.
- **Steady-state memory** — `--gpu-memory-utilization` (`LLM_GPU_MEM_UTIL=0.50`). Unified memory means there is no private VRAM: the 0.9 default would reserve 108 GB of the shared 128 GB and starve the embedder, reranker, Qdrant and the OS. Explicitly bounding it is correct practice here even though it was not the cause of the OOMs above.

Measured at `LLM_GPU_MEM_UTIL=0.50`: **KV cache 36.52 GiB = 398,832 tokens**.

`ninja` and `nvcc` must be on the unit's PATH or the engine dies with `FileNotFoundError: 'ninja'` after loading all 4 shards — which looks like a slow model load, not a missing build tool.

Verified

- `/v1/models` advertises `qwen3-vl-30b-a3b`.
- Coherent generation ("pong"), explicitly not the !!!!! NVFP4 failure mode.
- Tool calling via `--enable-auto-tool-choice --tool-call-parser hermes:` returns `get_weather` with parsed `{"city": "Brno"}`.
- Cold init engine (profile, create kv cache, warmup model) took **813.83 s** (torch.compile only 9.52 s of that — the rest is FlashInfer JIT at -j4).

NOT yet measured

Prefill/decode tok/s and real memory bandwidth are still unmeasured; ADR-0013's performance numbers remain inferred. The FP8 fallback (Qwen/Qwen3-VL-30B-A3B-Instruct-FP8) was never needed — it exists and is ungated if this checkpoint ever misbehaves.

Chapter 12b Model record — embedder

Embedding Model Record

Field	Value
Model	BAAI/bge-m3
Revision	5617a9f61b028005a4858fdac845db406aefb181 (resolved on the Spark, 2026-08-03; repo lastModified 2024-07-03)
Backend	FlagEmbedding 1.4.0 (BGEM3FlagModel), CUDA (devices="cuda:0"), FP16 (use_fp16=True) — ADR-0013 §4
Weights	models/embeddings/bge-m3/ — <code>pytorch_model.bin</code> (2.2 GB) + <code>sparse_linear.pt</code> head (load-bearing for sparse); path in <code>.env</code> <code>EMBED_MODEL_PATH</code>
Output	dense 1024-dim (L2-normalised) + sparse (lexical weights)
Languages	100+ incl. Czech and English
Max tokens	8192
Service	<code>services/embedder/server.py</code> — POST <code>http://127.0.0.1:8090/embed</code> (embeddings only)
Unit	<code>manifests/ragfarm-embedder.service</code> (host, CUDA)
Updated	2026-08-03

Weight format. This repo ships **no** `model.safetensors` — `pytorch_model.bin` is the only weight file BAAI publishes, so `fetch-encoder.sh`'s "fastest format" preference falls through to the pickle. That is fine on torch 2.13 (`weights_only` defaults to True); the old `safetensors-only` rule was a workaround for the abandoned NPU venv's very old torch and is retired. An earlier version of this record claimed `model.safetensors` (~2.3 GB) — that was never true of this repo.

Revision vs. the tested baseline. `docs/deployment.md` → "Tested model versions" records `50f9396f75618b3389c1fd1068a1ff58dc7b5b26` as the known-good baseline. Step 03's standing policy is to fetch **latest**, which resolved to the hash above. If retrieval quality ever looks wrong, re-fetch the baseline rev and compare before suspecting the pipeline.

The sibling cross-encoder reranker is recorded separately in `../reranker/MODEL.md` (it moved out of the embedder to its own GPU service, ADR-0008); the generative LLM in `../llm/MODEL.md`.

Chapter 12c Model record — reranker

Reranker Model Record (ADR-0008)

Field	Value
Model	BAAI/bge-reranker-v2-m3
Revision	latest (not pinned; fetched by <code>scripts/fetch-encoder.sh</code>)
Backend	<code>llama.cpp --reranking</code> , Vulkan / iGPU (Radeon 890M, RADV GFX1150), f16 GGUF
Weights	<code>models/reranker/bge-reranker-v2-m3/bge-reranker-v2-m3-f16.gguf</code> (1.15 GB, gitignored); path in <code>.env</code> <code>RERANK_MODEL_PATH</code>
HF source	BAAI/bge-reranker-v2-m3 (latest) — downloaded + converted to GGUF by <code>scripts/fetch-encoder.sh</code>
Output	one relevance score per (query, document). <code>llama.cpp</code> returns the raw logit rag-retrieval applies <code>sigmoid</code> → [0,1] (identical to <code>FlagReranker normalize=True</code>)
Role	cross-encoder rerank of the fused RRF candidate pool in <code>search_corpus</code>
Languages	100+ incl. Czech and English (XLM-RoBERTa-large family, sibling of bge-m3)
Latency	1.7 s / 40 candidates on the iGPU (was 36 s on CPU inside the embedder)
Service	dedicated <code>llama.cpp</code> server — POST <code>http://127.0.0.1:8081/reranking</code>
Unit	<code>manifests/ragfarm-reranker.service</code> (host, iGPU/Vulkan)
Updated	2026-07-21

Why a second llama.cpp server (not an embedder sub-endpoint)

The reranker shares nothing with the embedder anymore — different model, different device (iGPU vs CPU), different purpose. It is a plain `llama-server --reranking` instance, so there is no first-party code and no `services/reranker/` directory: the unit file is the whole deliverable. See ADR-0008.

Regenerating the GGUF (weights are gitignored)

The download + f16 GGUF conversion live in the fetch script — one command:

```
scripts/fetch-encoder.sh          # fetches bge-m3 + converts this reranker to  
GGUF  
scripts/fetch-encoder.sh --force  # re-fetch + re-convert  
scripts/fetch-encoder.sh --list   # other embedder+reranker pairs
```

It writes RERANK_MODEL_PATH into .env, which ragfarm-reranker.service reads.

Chapter 13 Instrumentation and tracing

Instrumentation and tracing — what these tools are, and what they

taught us **Status: a framework, not a finished toolkit**. These tools were built in July 2026 against the llama.cpp/AMD deployment to answer one question — where does the time and the context actually go? They answered it, the answers shaped the architecture, and then the ground moved under them: ADR-0013 replaced the generation engine and the Thinking models arrived. Treat them as instruments that need recalibrating, not as documentation of the current system.

This chapter replaces seven separate guides written in July. They documented an engine we no longer run, ports that never existed, and a set of "questions for David" that were answered a month ago. What was worth keeping is here; the rest is in git history.

What the tracing work established

These findings outlived the tools that produced them, and several are load-bearing in the current design:

- **Retrieval is dominated by the reranker.** The cross-encoder is 297 ms of a query, far more than Qdrant or the embedder. That is the cost of precision and it is worth paying — but it is why the reranker runs at lowered scheduling priority, so the interactive LLM wins contention.
- **Context growth per turn is the thing to watch, not context size.** Above roughly 150 tokens of growth per turn you are heading for overflow regardless of how much headroom you started with. This is what motivated bounding context by evicting old tool-result bodies rather than summarising, in `scripts/agent.py`.
- **Open WebUI's context loop is not observable from inside it.** Its compaction fires unpredictably and, after a summarisation pass, drops tool schemas — so a mutating action can narrate a success it never performed. That single finding is why `scripts/agent.py` exists at all: a loop we own, with the tool schemas permanently present so they cannot be compacted away.
- **The `:8000/rag` versus `:8104` distinction is a recurring trap.** Tools must talk to mcpo on `:8000`, not to the MCP server directly on `:8104`. Two separate tools were written against the wrong one.

What is wrong with them today

Deliberately not chased, because a rewrite would supersede it:

- **No concept of thinking models.** They cannot separate reasoning tokens from answer tokens. Against Qwen3-VL-Thinking, which routinely spends two thirds of its budget on reasoning, every timing they report is meaningless. This is the single reason not to trust their numbers today.
- **Docstrings still show the old port map** — localhost:8001/8002/8003. Those ports never existed in this deployment. The code was fixed on 2026-08-03 to resolve endpoints through `ragfarm_env.py`; the copy-pasteable examples in the docstrings were not. Trust the resolver, not the comments.
- `ragfarm_http_tracer.py` is proxy-based and needs to sit in the request path. It still holds bare `host:port` defaults.
- `chat_execution_tracer.py`'s **proxy mode is dangerous**. A previous session wired it into Open WebUI's database and **broke the model presets**. Do not re-plumb it without a plan. The retired `scripts/owui_trace_proxy.py` is the headstone.

Endpoint resolution

One resolver, used by every tool. Shell environment wins over `.env`, `.env` wins over built-in defaults, and the built-in defaults are the real deployed ports — so the tools work on a clean checkout with no configuration.

```
.venv/bin/python tests/tracing/ragfarm_env.py    # print resolved endpoints
```

The real port map is no longer duplicated here. `scripts/stack.sh list` prints it from the one place it is defined, and `scripts/stack.sh status` tells you which of them are actually answering. See `man docs/man1/stack.1`.

The tools

script	what it answers
<code>ragfarm_bench.py</code>	baseline tokens/s, prefill and decode, per prompt
<code>ragfarm_bench_extended.py</code>	the same with CSV/JSON export and prompt files
<code>ragfarm_bench_chatid.py</code>	context growth per turn, correlated by chat id
<code>ragfarm_rag_tracer.py</code>	candidate-pool evolution: Qdrant → RRF → reranker, per-stage tokens and latency
<code>ragfarm_integrated_tracer.py</code>	engine telemetry plus pipeline trace in one report
<code>ragfarm_tracer_simple.py</code>	minimal single-query timing
<code>chat_execution_tracer.py</code>	replay a canned session; <code>--demo</code> needs no LLM

ragfarm_http_tracer.py	proxy-based HTTP capture (see the warning above)
ragfarm_env.py	the endpoint resolver every other tool imports

```
.venv/bin/python tests/tracing/chat_execution_tracer.py --demo # no LLM
required
```

What superseded them

Most of what these tools were used for now has a purpose-built replacement, and that is the honest reason the rewrite has not been urgent:

- **Answer quality over time** → `scripts/test_regressions.py`, which replays `docs/prompts.md` against the live slot and judges the answers. See `docs/regression-testing.md`.
- **Model comparison** → `scripts/bench_ab.py`, and the raw results kept under `docs/measurements/`.
- **Service health** → `scripts/stack.sh status`, which probes all thirteen services and includes depth checks where a 200 can lie.
- **Retrieval quality on one query** → `scripts/rag_pool_inspect.py`.

Requirements for the rewrite

Whoever picks this up should carry forward:

1. **Thinking-model awareness.** Separate reasoning tokens from answer tokens and report both, or the timings mean nothing.
2. **Resolve endpoints through** `ragfarm_env`, never hardcode.
3. **Surface the ADR-0010 gate.** `search_corpus` already returns `_timing_ms.gate` with `floor_drop`, `kneedle_cut` and `kneedle_d` — the most useful retrieval diagnostic we have, and nothing reads it yet.
4. **Do not require sitting in the request path.** The proxy approach cost a broken Open WebUI database once already.

Chapter 14 Build progress (PROGRESS.md)

PROGRESS — blocker channel between the agent and Dave

This is the only file Dave writes into to steer the build. The agent appends **BLOCKED:** entries when it needs something only Dave can supply; Dave flips them to **UNBLOCKED:** once supplied. Linear build progress lives in `BUILD_STATE.md`, not here — this file carries only blockers and their resolution.

See `CLAUDE.md` Chapter 2 for exactly when and how this file is read and written. All timestamps are UTC. Newest entries are appended at the end. Resolved entries are kept as historical record, never deleted.

Format

```
BLOCKED: <NN-stepname> – <UTC timestamp>
  need:  <exactly what Dave must supply>
  where: <exact path / command / .env key / BIOS field involved>
  detail: <one or two lines of context>
```

Dave clears a blocker by editing that block's first line in place:

```
UNBLOCKED: <NN-stepname> – <UTC timestamp Dave cleared it>
  supplied: <what he did – file in place, creds in .env, toggle set, etc.>
  need:    <original need, left for the record>
  where:   <original where, left for the record>
  detail:  <original detail, left for the record>
```

Entries

```
UNBLOCKED: 01-npu-bringup — 2026-06-15T12:55Z need: Reboot the machine so the staged kernel parameter takes effect. where: /etc/default/grub — GRUB_CMDLINE_LINUX_DEFAULT now contains amd_iommu=force_isolation; update-grub has already been run; GRUB config is current. detail: The NPU PCI device (0000:c6:00.1) is in an IOMMU identity domain because the BIOS ACPI IVRS table has a unity-mapping entry for it. AMD IOMMU SVA (required by the amdxdna driver for PASID-based DMA) only works when the device is in a DMA-translated domain. The in-tree amdxdna driver logs "SVA bind device failed, ret -95" on every open(). The fix is to add amd_iommu=force_isolation which overrides IVRS unity-mapping entries and forces all devices into isolated (translated) domains. After reboot: re-run xrt-smi examine and python infra/npu/quicktest.py to verify the gate (NPU Strix in examine output, "Test Finished" from quicktest).
```

```
BLOCKED: 05-mcp-placement — 2026-06-18T11:30Z need: live OpenNebula access — ONE_XMLRPC endpoint + ONE_AUTH credentials, and network reachability to the ON frontend where: .env keys ONE_XMLRPC, ONE_AUTH (shape in .env.example) detail: PoC MiniPC is not yet placed on the live network; ON
```

access lands at deployment. The placement MCP code + XML parsing are unit-tested; only the live one.vm.info/one.vmpool.info round-trip is unverified.

BLOCKED: 06-mcp-fs-host-control — 2026-06-18T11:30Z need: live OpenNebula access (same as 05) before host-control real actions where: .env keys ONE_XMLRPC, ONE_AUTH detail: fs-agent (sandboxed read) can be implemented and tested now, but host-control's drain-then-reboot is via ON and cannot be verified without a live cluster. Keep host-control dry-run/confirmed; do not enable real actions until ON is reachable at deployment.

NOTE: ADR-0006 activated — 2026-07-14 UTC owner: Dave Kubicek (owner-directed capability change, not a build-step) summary: Content-addressed corpus sync with SQLite manifest + alias switch + autonomous watcher deployed and validated end-to-end. changes: - Step-04 "ingester frozen" rule superseded by ADR-0006 for owner-driven work; xlsx_tables.py parser remains off-limits to build agents. - Project venv moved: .venv on python3.12 replaces /opt/ryzenai/venv for both ingester and embedder. NPU env stripped from ragfarm-embedder.service. - ragfarm-ingester-watcher.service installed, enabled, wired into ragfarm-stack. - Legacy physical corpus collection migrated to alias corpus -> corpus_20260714-025915_901856e8 via -recreate (zero blackout after migration). - infra/compose.yaml ingester: block removed (superseded by host watcher). gate: All \$9 gates passed; see logs/ingester-adr0006.log.

NOTE: OWUI serving tuning (context handling + tool-calling) — 2026-07-14 UTC owner: Dave Kubicek (owner-directed debugging, not a build-step) problem: Multi-turn RAG chats in Open WebUI hit a slow CPU-busy "before-tool" delay, context overflow (hard 400 errors), stalls, and degraded tool-calling (model narrated instead of calling reboot_host). Root cause: unbounded accumulation of REPLAYED tool-result blobs in OWUI conversation history — each turn resends the full prior tool outputs (a single docx "How-Tos" chunk is 4k tokens), so prompts reached 10-14k+ tokens. A warm llama prompt-cache had hidden the cost for days; a llama restart exposed it (cold full re-prefill of the whole history). NOT caused by the .venv/embedder move or the watcher. changes: - manifests/ragfarm-llama.service: -c 16384 -> -c 32768; added -context-shift -keep 3072. NOTE: -context-shift only rescues GENERATION overflow, NOT an oversized prompt (verified: 28k-token prompt still returns 400 exceed_context_size_error). So 32k only raises the ceiling; it does not make overflow impossible. -keep 3072 preserves system prompt + tool schemas across shifts. (-parallel 1 was tried then reverted to default 4 slots; prompt-cache restore makes single-slot safe but 4 slots matched the proven baseline.) - OWUI model "ragfarm": function_calling native -> default.

THIS was the effective fix — restored immediate/correct tool calls, faster generation, and cut per-turn accumulation. Stored in the openwebui_data docker volume (webui.db), NOT version-controlled; revert with a one-field DB update. tradeoff / PROD caveat: - In "default" mode the model answers a REPEATED question from context instead of re-calling the tool (observed: 2nd "Kde bezi sftp-gw?" and 2nd "Jak se prihlasim?" made zero tool calls). Harmless in test, but STALE ANSWERS are unacceptable for production. - PROD PLAN: return to function_calling = "native" for the real deployment, paired with a larger/newer model + bigger context window (which is what lets native mode carry the accumulated trace without overflowing). - A reboot sentence that prefaced one login answer was verified as a pure context echo (that turn called zero tools) — no host-control execution. deferred lever #2 (accumulation suppression via retrieval; NOT applied): - In services/rag-retrieval/server.py, cap each result's "text" (800 chars) and lower default k (5 -> 4) so each tool result is 1k instead of 4k tokens, slowing history growth. Parked because it degrades answer grounding to patch a plumbing issue; the clean fix is finer docx chunking in the (frozen) ingester parser. Keep #2 as a fallback if context buildup returns.

UNBLOCKED: 02-vllm-serving — 2026-08-03T13:05Z supplied: Dave's decision, given in-session: "Absolutely, give vLLM its own .venv-vllm." Step 02 installs vLLM into a DEDICATED venv at .venv-vllm; .venv stays the pinned CUDA-13 environment for the embedder, the frozen ingester and the MCP services. (original blocker below, kept as the historical record) BLOCKED: 02-vllm-serving — 2026-08-03T12:34Z need: a decision on WHERE vLLM is installed. BUILD_STATE step 02 command 1 says "vLLM into the step-01 venv" (.venv/bin/pip install -U vllm). Verified by dry-run on the Spark: doing that dismantles the environment step 01 just gated. Recommendation: give vLLM its OWN venv (.venv-vllm) and leave .venv as the pinned CUDA-13 env for embedder/ingester/MCP. where: BUILD_STATE.md step 02 command 1; scripts/deploy.sh phase_venv; manifests/ragfarm-vllm.service (not yet written); .venv vs .venv-vllm detail: .venv/bin/pip install --dry-run -U 'vllm>=0.22.0' resolves vllm 0.26.0 and would change, in the SAME venv the ingester and embedder run from: torch 2.13.0+cu130 -> 2.11.0 (plain PyPI, NOT a cu130 build — this alone undoes step 01 gate 3, sm_121) numpy 1.26.4 -> 2.3.5 (major) transformers 4.57.6 -> 5.14.1 (major) huggingface_hub 0.36.2 -> 1.26.0 (major) nvidia-cudnn-cu13 9.20.0.48 -> 9.19.0.56 nvidia-nccl-cu13 2.29.7 -> 2.28.9 protobuf 7.35.1 -> 6.33.6 fastapi 0.137.1 -> 0.136.3 numpy 1.x->2.x and transformers 4.x->5.x sit directly under the FROZEN services/ingester parser and FlagEmbedding/BGE-M3 (step 03/04), so this is not just a torch problem. Nothing in .venv needs to import vllm: vLLM is reached over HTTP on :8080 and ADR-0003 keeps the retrieval pipeline serving-

engine agnostic, so a second venv costs only disk. Also for the record: resolved stable vLLM is 0.26.0, not the v0.22.x BUILD_STATE anticipates (" $\geq 0.22.0$ " is still satisfied, so PR #40082 is in), and flashinfer-python 0.6.14 IS in the resolved set, so the b12x native-FP4 target looks reachable. Step 01 remains DONE and intact — no package was installed; -dry-run only.

BLOCKED: 07-agent-wiring — 2026-08-04T10:40Z need: ONE browser click-through by Dave to close the RAG-only gate. Everything is wired and proven up to OWUI's own tool-execution loop, which an API caller structurally cannot trigger. where: `http://127.0.0.1:3000` (or LAN IP:3000), login `admin@ragfarm.local`, password in repo-root .env key `OWUI_PASSWORD` (generated, PLEASE ROTATE) detail: In the UI pick model "ragfarm-vision" and send these two prompts: 1. "Kolik vCPU a RAM ma VM hsmbvxi001ts?" (sparse exact-identifier) 2. "Jak se prihlasim do EPC? Popis postup." (czech semantic / dense) PASS = each answer is grounded in corpus content AND the turn shows a `search_corpus` tool call. Confirm the call really ran with: `docker logs infra-rag-retrieval -tail 20 | grep search_corpus` which prints one "`search_corpus q='...' -> n/40 cand`s" line per call.

WHY THIS NEEDS A HUMAN. OWUI resolves the tool and forwards the schema to vLLM correctly (verified in its DEBUG log:
`tools=[tool_search_corpus_post]`
with the full parameter schema). The model then emits a correct `tool_call`. But OWUI executes registered OpenAPI tool servers inside `process_chat_response`, which is driven by a chat/websocket session; an API request carries `chat_id=''` and `session_id=None`, so OWUI streams the `tool_call` back to the caller instead of executing and looping. That is an OWUI API-vs-UI asymmetry, not a defect in our wiring, and no amount of API-side probing can close it.

PROVEN ALREADY (logs/07-agent-wiring.log):

- mcpo mounts `/rag/search_corpus` (+ placement 2 ops, host-control 1)
- `search_corpus` via mcpo: hostname -> correct row @0.99; czech query -> correct EPC login chunk
- vLLM emits a correct `tool_call` for OWUI's exact tool schema
- closed-loop (`tool_call` -> mcpo -> feed back) grounded answers in BOTH modes: "VM hsmbvxi001ts ma 8 vCPU a 64 GB RAM", and a correct Czech EPC login procedure with real URLs/hosts from the corpus
- OWUI preset `ragfarm-vision` bound to live alias `qwen3-vl-30b-a3b` with `toolIds [server:0, server:1, reboot_guarded]`, `function_calling=native`

If the UI click passes, flip this to UNBLOCKED and step 07 is DONE (the deploy fragment is the only remaining work). If the UI ALSO fails to execute the tool, that is a real OWUI integration defect and I should look at the direct/client-side tool-server path next.

NOTE: otevřené položky po A/B testu — 2026-08-07 owner: Claude Opus 5 (zápis), rozhodnutí Dave

- 1) Porucha "narace tool callu místo jeho vyvolání". V jednom z 20 běhů MoE nevygeneroval volání nástroje a místo odpovědi vypsál vlastní úvahu ("`<tool_call>` tags. Let me double-check RULE 0..."). Dave rozhodl VYNECHAT ji z docs/vyzkumna-zprava-2026-08-07.md a nechat si ji jako téma pro příští weekly. Žádné číslo ve zprávě kvůli tomu nebylo upraveno — po opravě promptu se v dalších 60 bězích neopakovala, takže 9/10 a 10/10 jsou skutečně naměřené. Mitigace v RULE 1 ("EMIT THE CALL, DO NOT DESCRIBE IT") je nasazená, ale NEOVĚŘENÁ proti tomuto konkrétnímu selhání.
- 2) Statistická síla. Rozdíl MoE 9/10 vs. dense 10/10 je jediná odpověď a na deseti bězích neunesl silný závěr. 50 běhů by ukázalo, jestli je reálný. Nástroj: scripts/bench_ab.py `<out.json>` 50
- 3) Doporučení ze zprávy: MoE jako výchozí model, dense ve slotu 1 pro úlohy, kde je přesnost kritická. Zatím nerealizováno jako trvalá konfigurace.